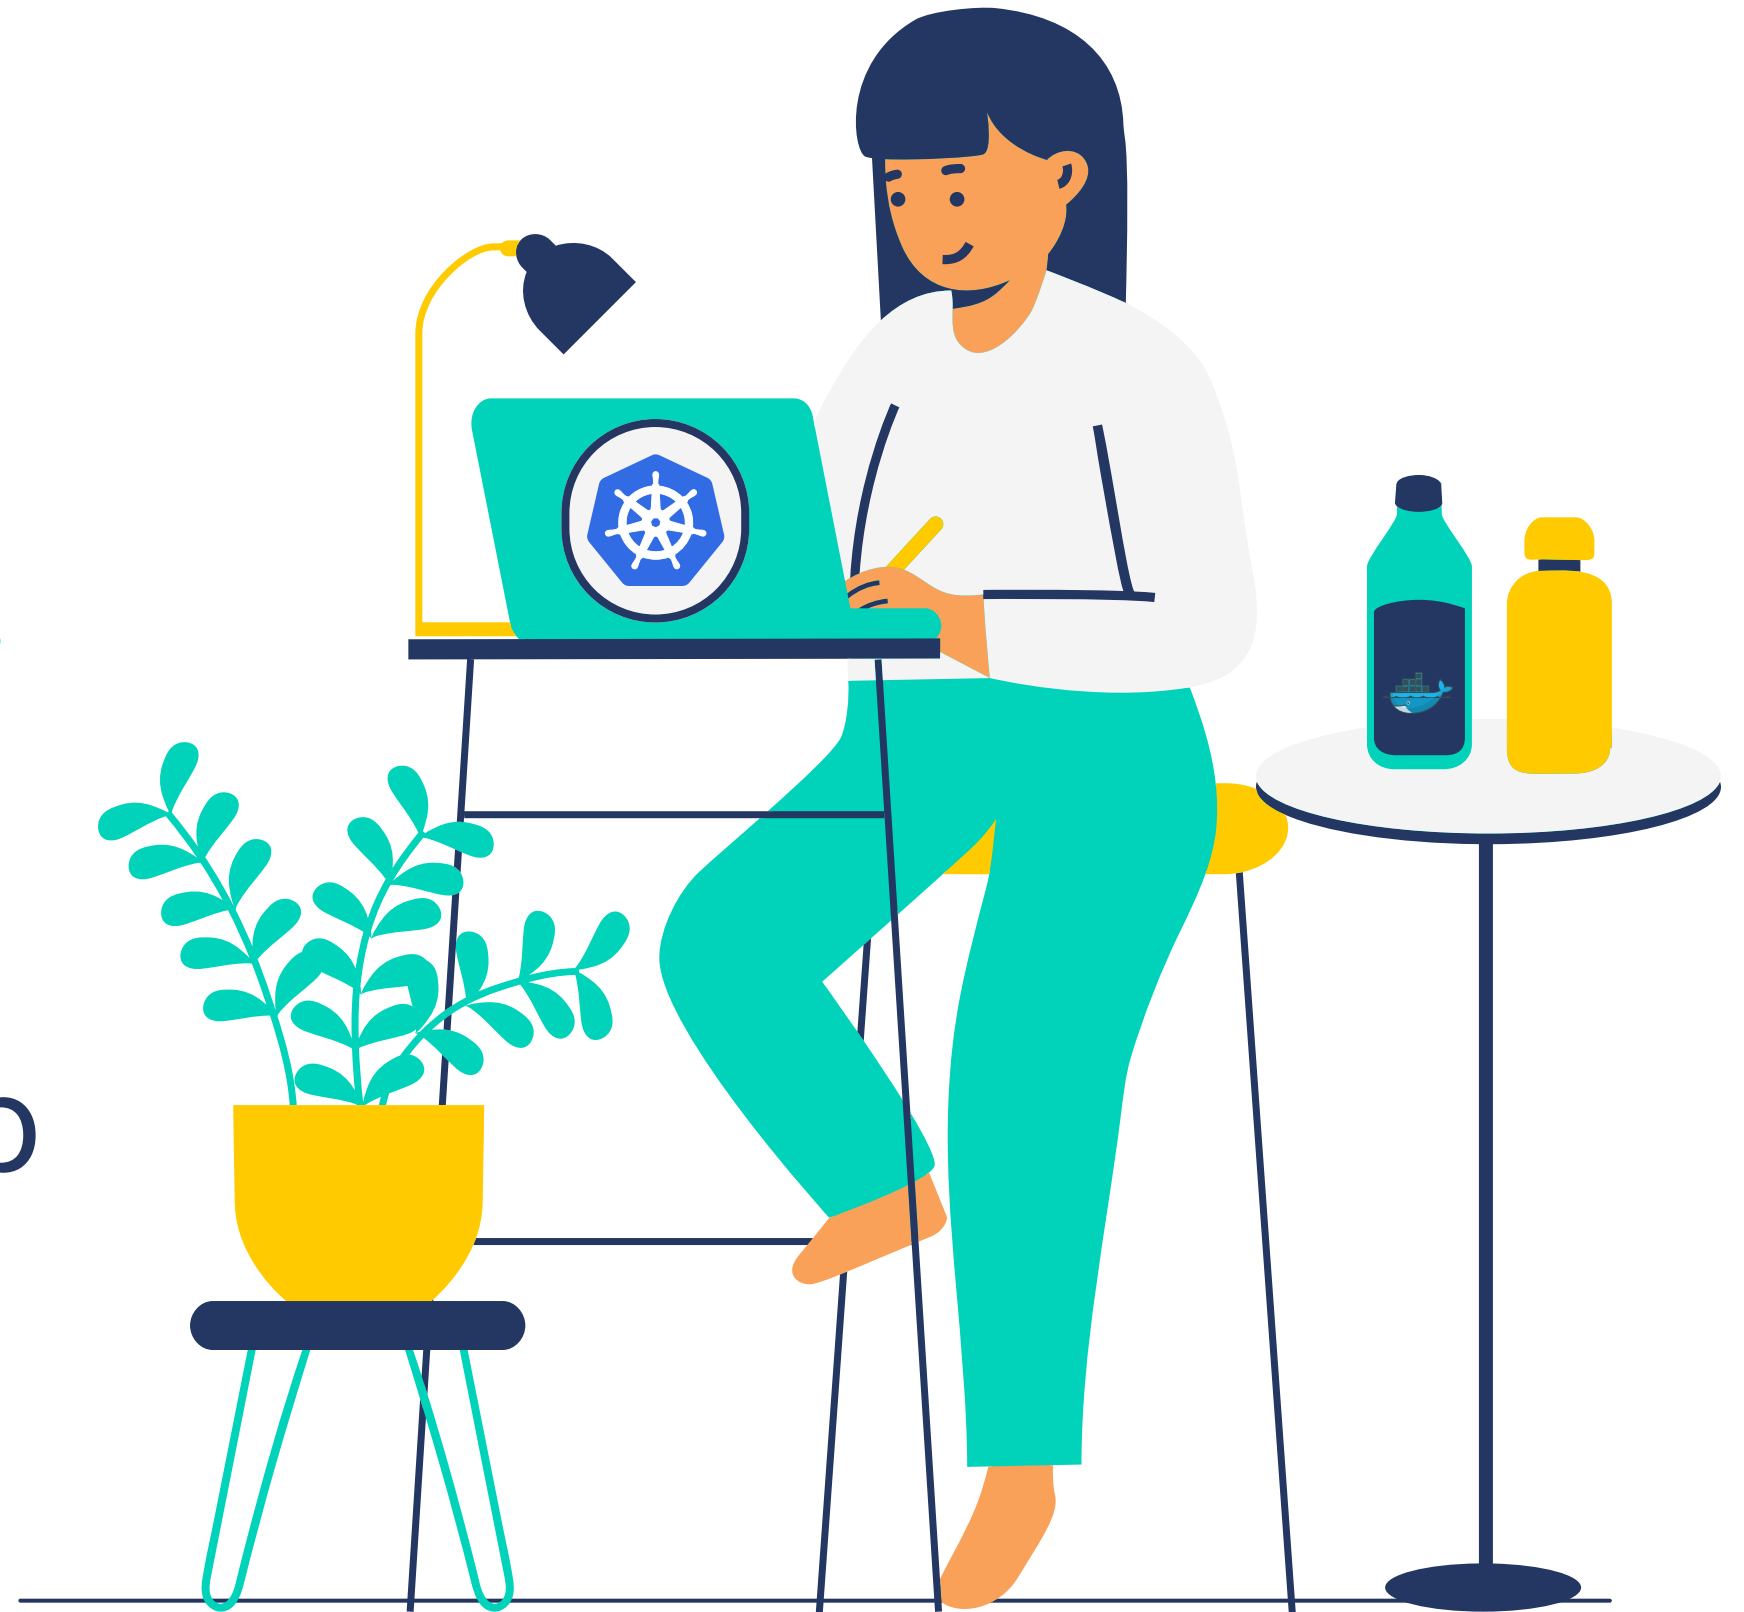


# Containers & Kubernetes

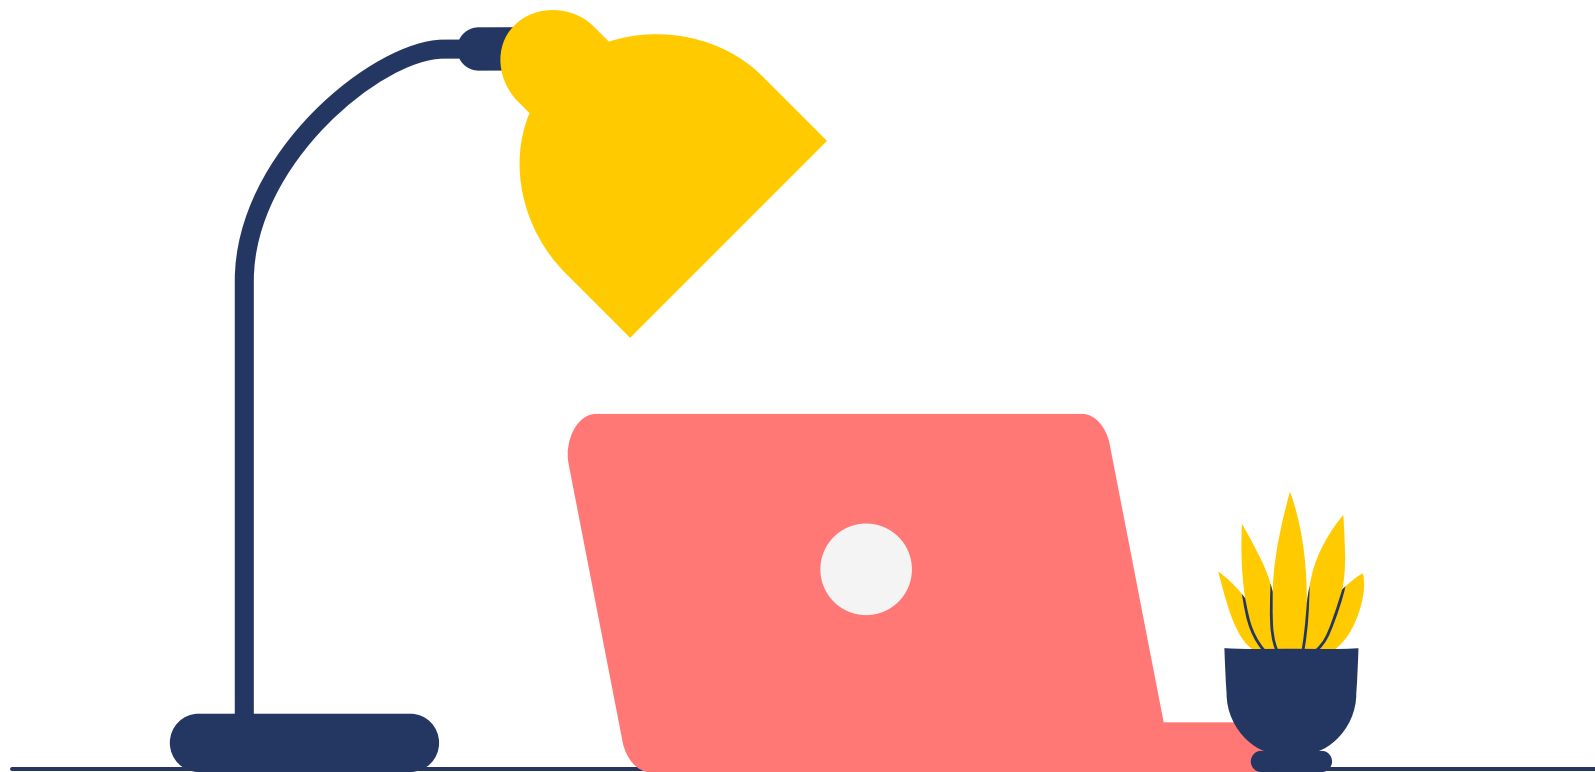
## A Beginner's Workshop

Rigerta Demiri & Katharina Sick  
#theNewITGirls | November 17, 2023



# About today

Let's do a step-by-step journey into the world of containers and Kubernetes.



## Containers & Docker

What are containers? Why are they cool?  
How to create & use them?



## Kubernetes

What is Kubernetes? Why do so many companies need it? How to get started?



## Troubleshooting and Questions

Feel free to ask any time! :)



# Hello, I'm Rigerta 🙋



Data Engineer @GitLab



Vienna Data Engineering Meetup



[linkedin.com/in/rigerta/](https://linkedin.com/in/rigerta/)



ÖSTERREICHISCHES  
ROTES KREUZ

Aus Liebe zum Men!



VIENNA DATA ENGINEERING  
**MEETUP**

# So, what is Docker?

An open-source tool used to turn your **application** into a **container** that can easily be deployed in any other system.





# First things first, what is an Application?

Any piece of **Software** that performs a specific function either for an end user or for another application.

An **Application** is written in a specific **Programming language** and typically has one or more **Dependencies**.

---



# Cool, so what is a Container?

A self-contained, runnable software application or service.

---

# And how do we do this with Docker?

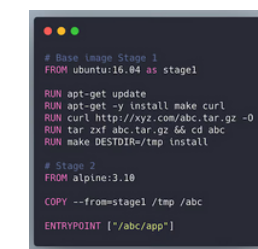
We write a **Dockerfile**, which is “a recipe” defining how to build a **Docker Image** and when we run the image so we can reach the application, we are in a **Container**



# Key Concepts



- **Application**
  - Any piece of software you write, to fulfil a goal for end users or other applications
- **Dependencies**
  - All software libraries an application needs to be able to run successfully
- **Dockerfile**
  - A text file containing a set of instructions
- **Docker image**
  - A read-only blueprint that includes container-creation instructions
  - An **executable** application artifact
- **Docker container**
  - A running instance of a Docker **image** that gets created when the **\$ docker run** command is implemented
  - Multiple containers can run from the same Docker image
- **Docker registry**
  - A storage and distribution system for container images (can be **private** or **public**)
  - Docker Hub ([hub.docker.com](https://hub.docker.com)) is the **official public** Docker registry



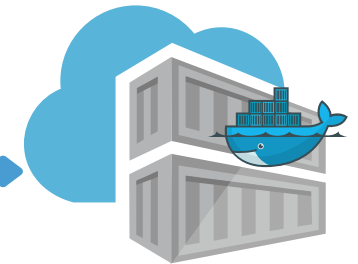
Dockerfile

— build —→



Docker Image

— run —→



Docker Container



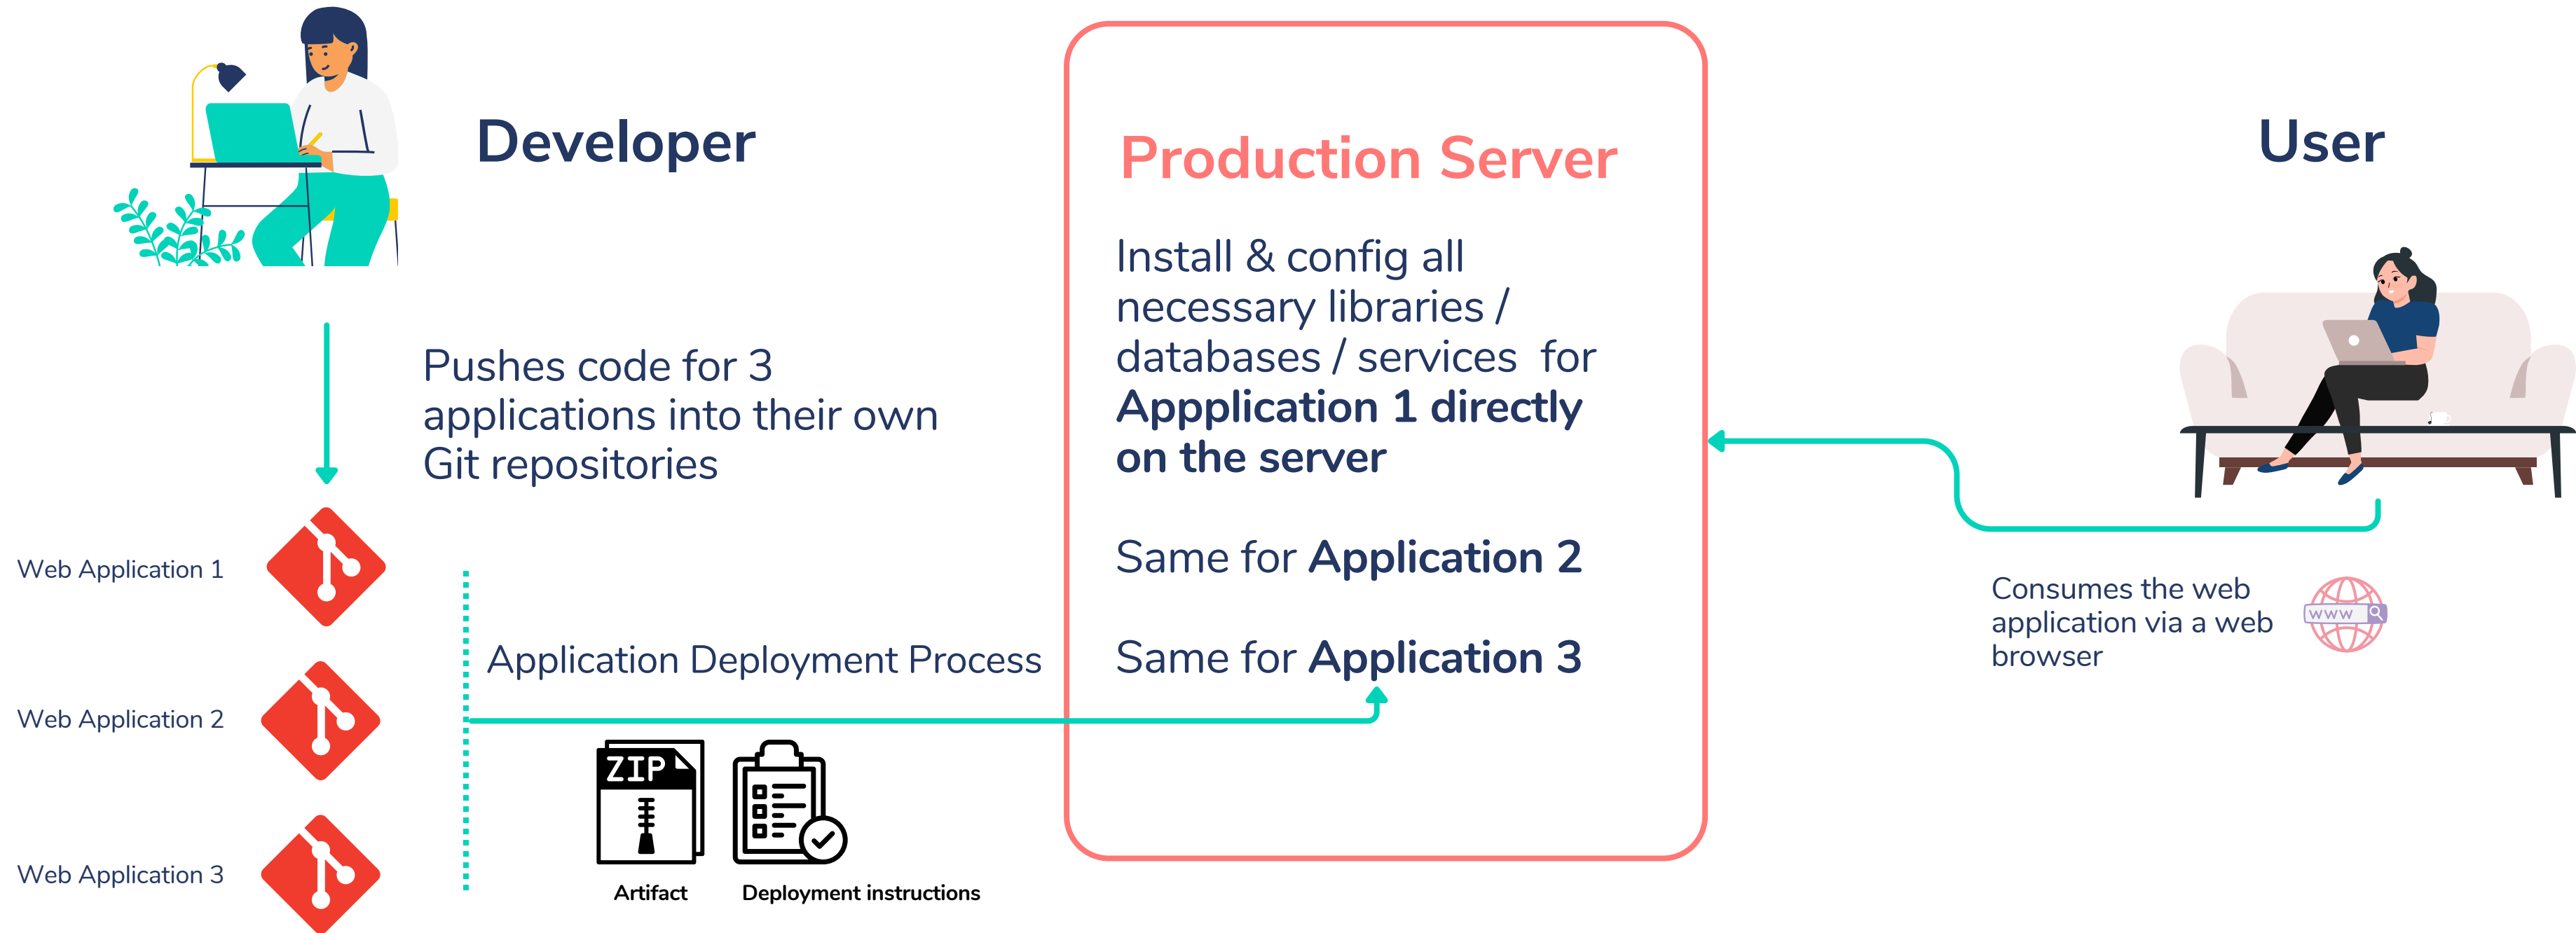


# OK, let's take a step back

Why do we even **need** Docker?  
How were applications deployed before it?

---

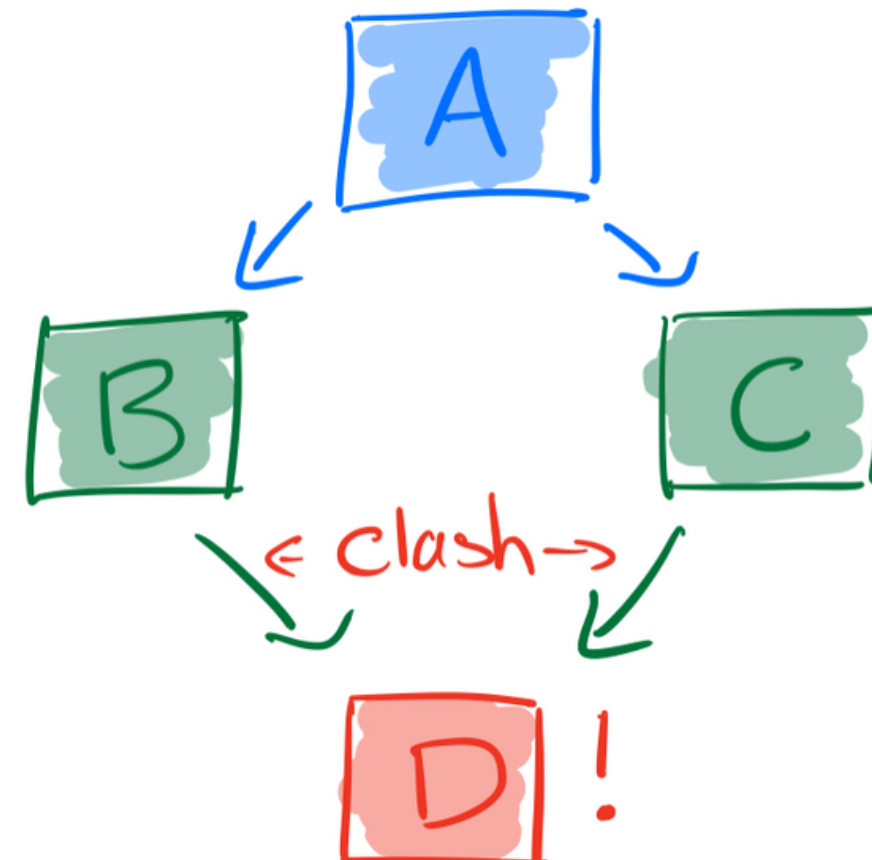
# Traditional Application Deployment



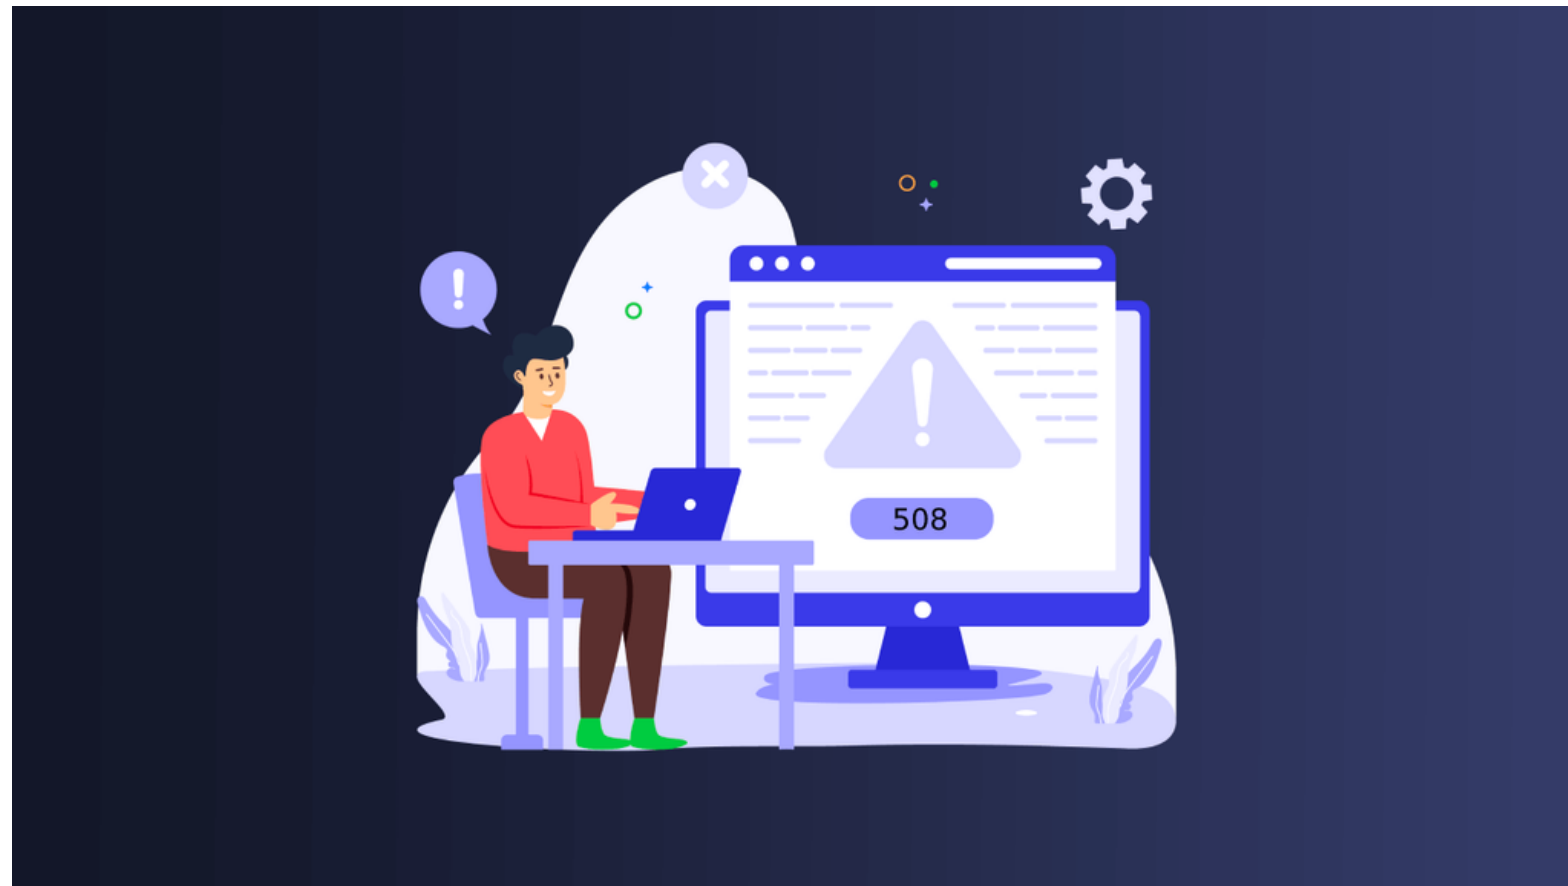
# What is the problem with that?

The **same** production server is usually used for **multiple** applications, for efficiency, **but**:

- Error prone process
- Different applications might need different
  - **software & software versions**
  - **libraries & library versions**
- This causes:
  - dependency issues
  - a painful **deployment process**



# What is the problem with that?

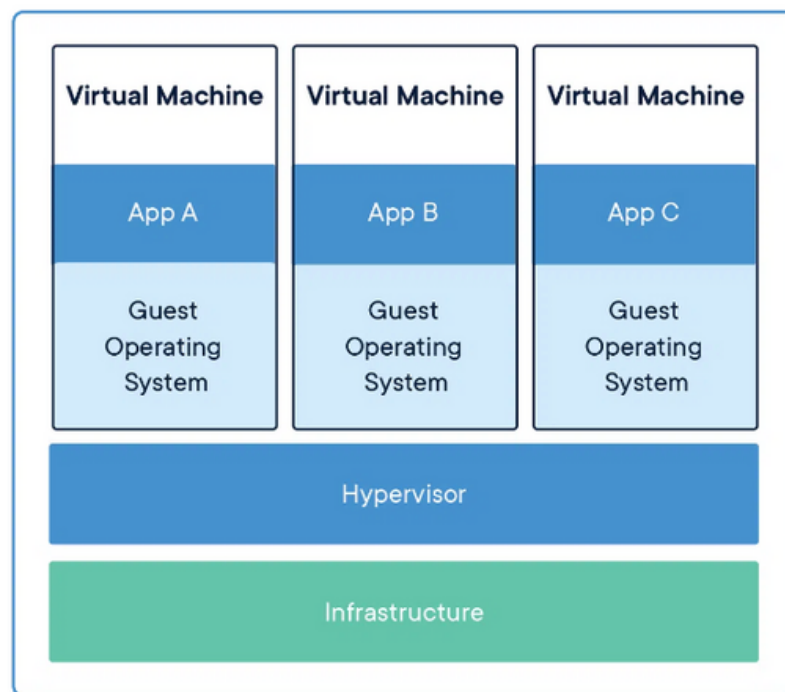


Memory and processing power of the production server are **limited**



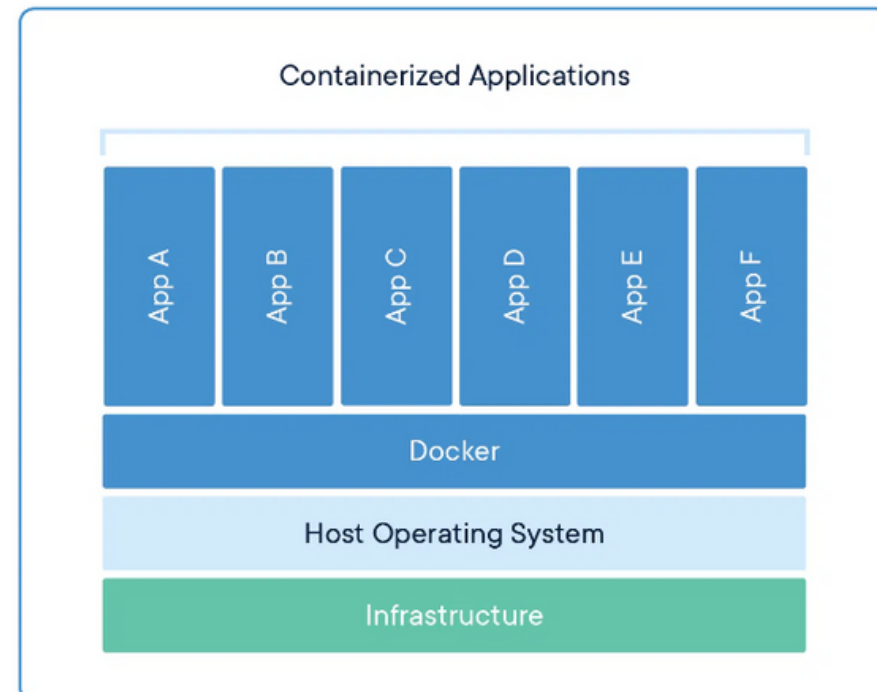
This is where Kubernetes & Docker play very well together and you will hear all about it in Part 2 of this workshop by Katharina

# Enter Docker



## VM Virtualization

OS Kernel & the application layer of the OS

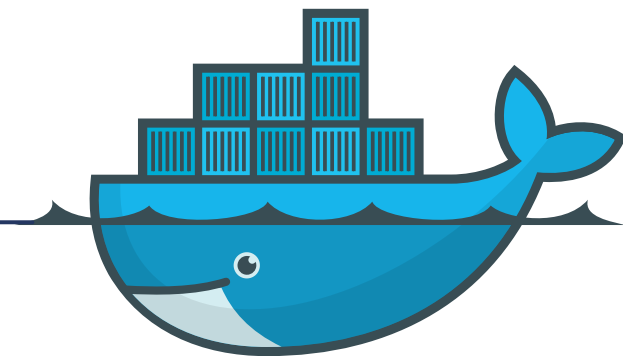


## Docker Virtualization

Only the application layer of the OS

## Benefits


- **Isolation**
  - No conflicting with the host system
- **Packaged application and all its dependencies** (code, libraries, system binaries etc)
  - No more “it works on my machine” issues
- **Reliable deployment**
  - A **self-contained** image is used to deploy, making the deployment independent of the OS or other details of the host system
  - Docker runtime is **the only** necessary installation



# Enter Docker

## More benefits

- **Better resource consumption management**
  - Configurable memory & cpu a container can use ([docs on resource constraints](#))
- **Efficiency**
  - A side effect of the lightweight, efficient isolation model of containers
  - **Many** Docker containers can run on a single production server



# Demo Time



- 💡 Step by step Docker introduction
- 💡 Docker Registry
- 💡 Use existing Docker Images
- 💡 Create your own Docker Image

# First steps



- **Docker running locally**

- Instructions to download and install Docker
- I run/prefer **Rancher Desktop**
  - An open-source application that provides all the essentials to work with containers and Kubernetes on the desktop



- **An IDE or text editor to use for editing files.**

- I prefer VSCode - [Download link](#)




- **Free Docker Account**

- You can sign-up for a free Docker account (<https://hub.docker.com>) and access free unlimited public repositories
- **Not mandatory**, images are available to unauthenticated users too



# Demo Part 1



- **Where to find Docker Images?**

- Docker Registries
  - Public (Docker Hub)
  - Private (AWS ECR, Google Container Registry, etc)



- **How to create a Container?**



- **How to access a containerised application**

- Port Binding



- **How to stop Containers**



- **How to remove Docker Images and Containers**

# Demo Part 2

## Creating your own Docker Image



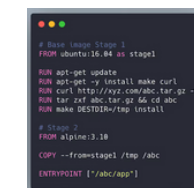
- Creating a Streamlit Application



- Running it locally with `streamlit run`



- Building a Docker Image for this Application



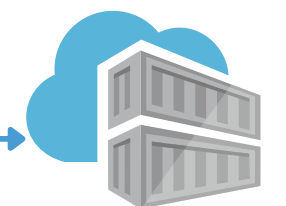
Dockerfile

— build —→



Docker Image

— run —→



Docker Container

# Main Takeaways



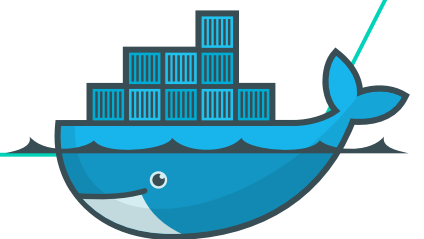
Docker is an open source virtualization technology



With Docker you can package your application into a single runnable artifact - Docker Image



Docker Images will run everywhere, independent of the platform you deploy it to



# Where to go from here?

## Explore the official documentation

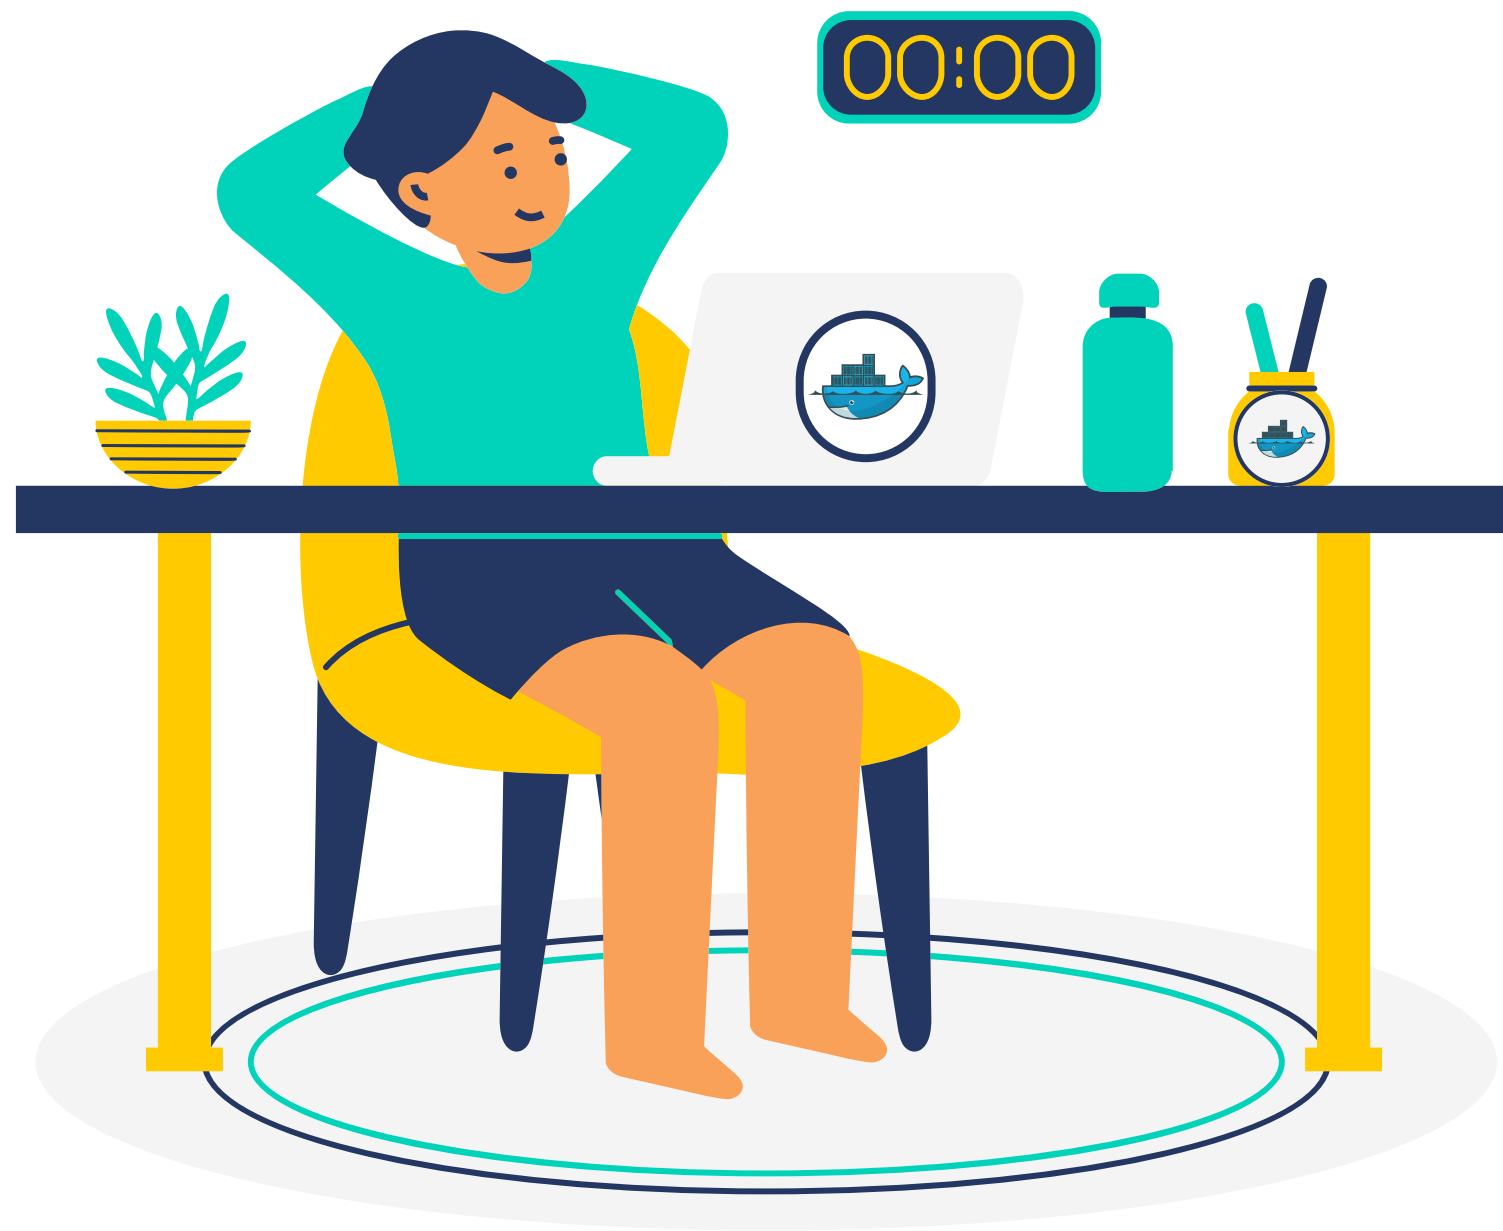
- <https://docs.docker.com/get-started/>
- <https://www.docker.com/resources/what-container/>
- <https://docker-curriculum.com/#docker-compose>
- <https://rancherdesktop.io/>
- <https://docs.streamlit.io/>

## Watch YouTube videos

- [Techworld with Nana](#)

## Experiment with side projects

- Build a simple application in any language of your choice and containerise it
- Play around and experiment with Docker
- Have fun!



Thank you!

Happy to take any  
questions.

# Kubernetes

What is Kubernetes? Why do so many companies need it? How to get started?



# What is Kubernetes?

## Container Orchestration System

- Takes care of all the containers you want to run

## Automates Deployment, Scaling and Management

- Deploys your applications based on your description (declarative configuration)
- Scale on demand
- Self-healing capabilities

## Open Source

- Open-sourced by Google in 2014
- Part of the Cloud Native ecosystem

**Kubernetes = Helmsman/Pilot = k8s**



# Why is it so useful?

It helps to organize and manage a huge amount of containers

It can scale your applications in busy times

It makes communication between various components easier

It allows you to update your applications without downtime

It has a very big and supportive community





# When to use it?

## **If you have a huge amount of applications and services running**

- Basically, it makes the most sense in enterprise environments

## **It can be an overkill in many situations and there are good alternatives for smaller projects**

- Hosting your own server
- Serverless functions (Google Cloud Run, AWS Lambda,...)



# Core Components



# Core Components

## Node

- A worker machine (physical or virtual)
- A Kubernetes cluster runs on one or multiple nodes



# Core Components

## Node

- A worker machine (physical or virtual)
- A Kubernetes cluster runs on one or multiple nodes

## Pod

- Runs one or multiple containers



# Core Components

## Node

- A worker machine (physical or virtual)
- A Kubernetes cluster runs on one or multiple nodes

## Pod

- Runs one or multiple containers

## Deployment

- Manages a set of pods and ensures they are up and running



# Core Components

## Node

- A worker machine (physical or virtual)
- A Kubernetes cluster runs on one or multiple nodes

## Pod

- Runs one or multiple containers

## Deployment

- Manages a set of pods and ensures they are up and running

## Service

- Exposes pods to other pods and services in the cluster



# Core Components

## Node

- A worker machine (physical or virtual)
- A Kubernetes cluster runs on one or multiple nodes

## Pod

- Runs one or multiple containers

## Deployment

- Manages a set of pods and ensures they are up and running

## Service

- Exposes pods to other pods and services in the cluster

## Namespace

- Isolates resources in a cluster



# Let's go!

## Install kind

- We'll use kind to locally run a Kubernetes cluster
- <https://kind.sigs.k8s.io/docs/user/quick-start/#installation>
- Verify if kind is working:
  - **kind version**

## Install kubectl

- We'll use kubectl to interact with the Kubernetes cluster
- <https://kubernetes.io/docs/tasks/tools/#kubectl>
- Verify if kubectl is working:
  - **kubectl version --client**



# Let's really go!

## Create a local Kubernetes cluster

- `kind create cluster --name kubernetes-playground`

## Execute first commands in the cluster

- `kubectl cluster-info`
- `kubectl explain pod`

## Take a look at kubeconfig

- `cat ~/.kube/config`

## Explore the newly created cluster

- `kubectl get namespaces`
- `kubectl get nodes -o wide`
- `kubectl get pods`

# Let's deploy something !

Create a file called deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-kubernetes
spec:
  template:
    metadata:
      labels:
        app: hello-kubernetes
    spec:
      containers:
        - image: paulbouwer/hello-kubernetes:1
          name: hello-kubernetes
          ports:
            - containerPort: 8080
              name: http
      selector:
        matchLabels:
          app: hello-kubernetes
```

# Let's deploy something II

## Create a namespace

- `kubectl create namespace hello-kubernetes`

## Create a deployment

- `kubectl apply -n hello-kubernetes -f deployment.yaml`

## Take a look at the created resources

- `kubectl -n hello-kubernetes get deployments`
- `kubectl -n hello-kubernetes get pods`
- `kubectl port-forward -n hello-kubernetes deployment/hello-kubernetes 8080:8080`

# Let's set environment variables

## Add an environment variable

- `spec.template.spec.containers[0].env:`
  - - `name: MESSAGE`
  - `value: 'Hello #theNewITGirls! :)'`

## We can also read information from Kubernetes

- `spec.template.spec.containers[0].env:`
  - - `name: KUBERNETES_NAMESPACE`
  - `valueFrom:`
  - `fieldRef:`
  - `fieldPath: metadata.namespace`
  - - `name: KUBERNETES_POD_NAME`
  - `valueFrom:`
  - `fieldRef:`
  - `fieldPath: metadata.name`
  - - `name: KUBERNETES_NODE_NAME`
  - `valueFrom:`
  - `fieldRef:`
  - `fieldPath: spec.nodeName`

# Let's use a configmap

Create a file called configmap.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: hello-kubernetes-config
data:
  message: "Hello from the config map! :)"
```

Apply it

- `kubectl apply -n hello-kubernetes -f configmap.yaml.yaml`

Reference it

- `spec.template.spec.containers[0].env:`
  - - `name: MESSAGE`
  - `valueFrom:`
  - `configMapKeyRef:`
  - `name: hello-kubernetes-config`
  - `key: message`

# Let's use a secret

Create a file called secret.yaml

```
apiVersion: v1
kind: Secret
metadata:
  name: hello-kubernetes-secrets
stringData:
  message: "I am a secret value!"
```

Apply it

- `kubectl apply -n hello-kubernetes -f secret.yaml`

Reference it

- `spec.template.spec.containers[0].env:`
  - - `name: MESSAGE`
  - `valueFrom:`
  - `secretKeyRef:`
  - `name: hello-kubernetes-secrets`
  - `key: message`

# Let's scale it

## Set replicas

- `spec.replicas: 3`

## Take a look at the created resources

- `kubectl -n hello-kubernetes get deployments`
- `kubectl -n hello-kubernetes get pods`
- `kubectl port-forward -n hello-kubernetes deployment/hello-kubernetes 8080:8080`

# Let's update without downtime

## Change the update strategy

- `spec.strategy.type: RollingUpdate`

## Take a look at the created resources

- `kubectl -n hello-kubernetes get deployments`
- `kubectl -n hello-kubernetes get pods`
- `kubectl port-forward -n hello-kubernetes deployment/hello-kubernetes 8080:8080`



# Let's create a service

Create a file called service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: hello-kubernetes
spec:
  selector:
    app: hello-kubernetes
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 8080
```

Create a service

- `kubectl apply -n hello-kubernetes -f service.yaml`

# Let's clean up

## Delete the Kubernetes cluster

- `kind delete cluster --name kubernetes-playground`

# Where to go from here?

## Explore the official documentation

- <https://kubernetes.io/docs/home/>

## Watch YouTube videos

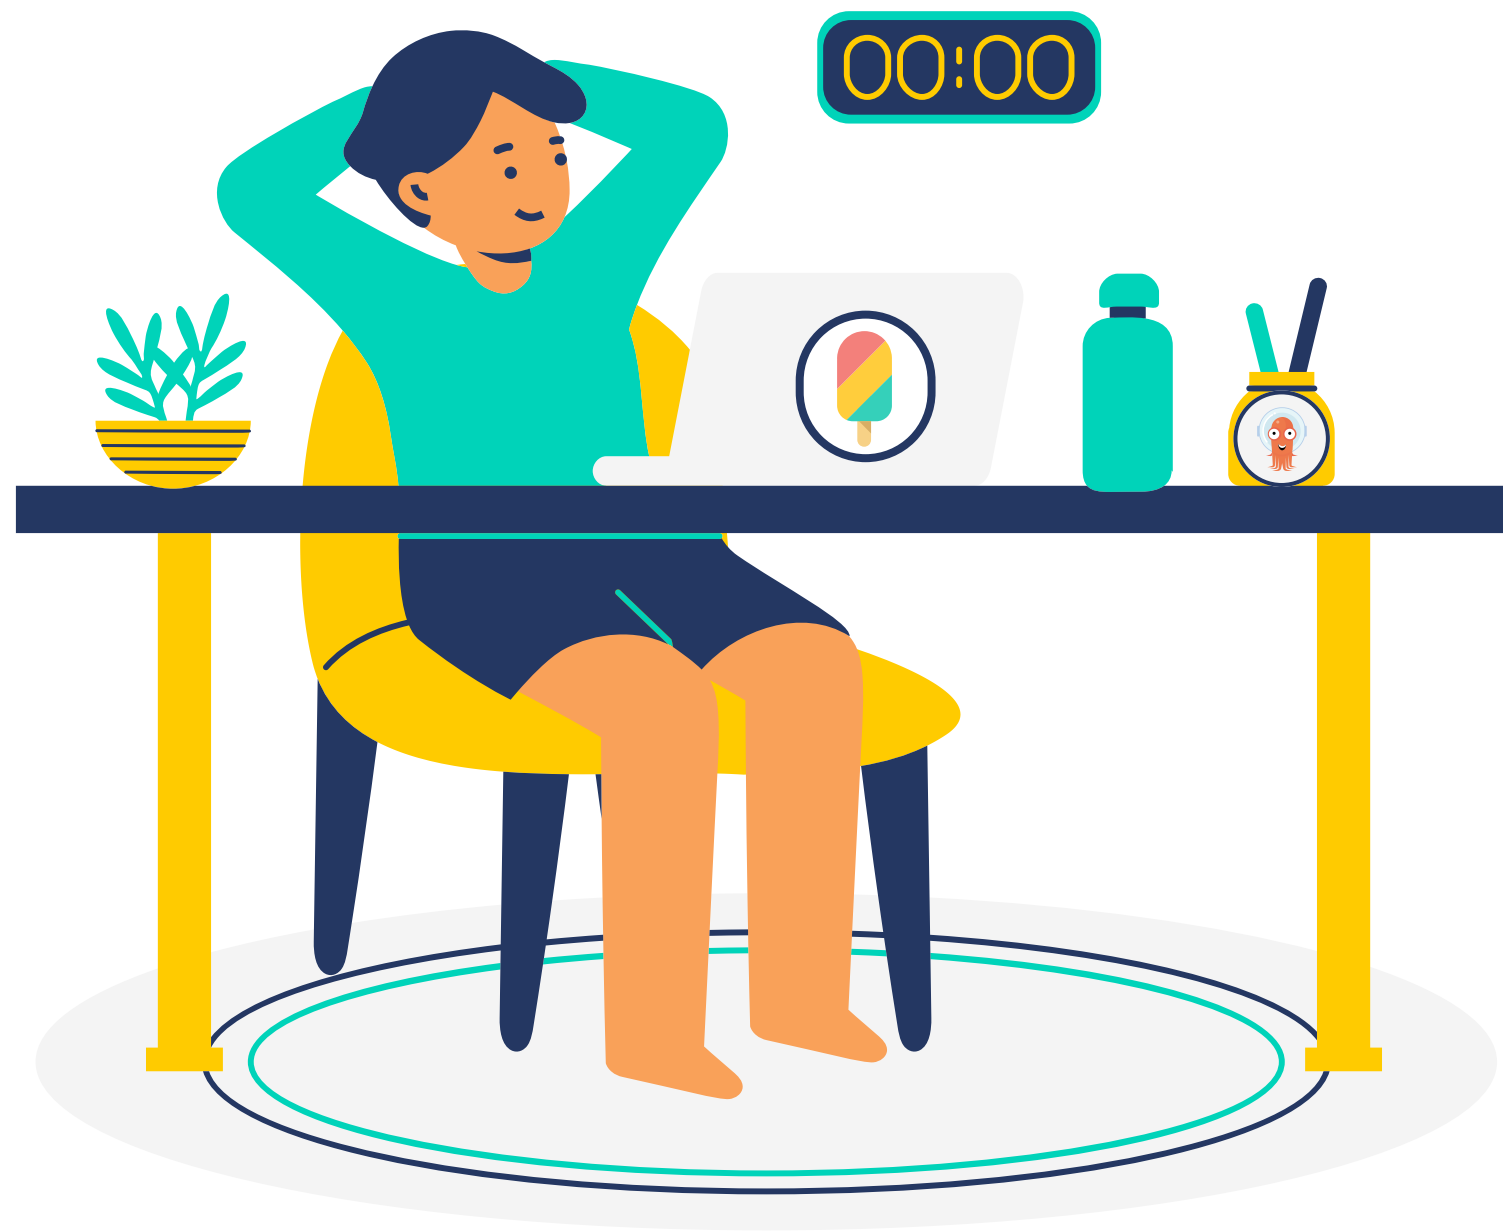
- [Techworld with Nana](#)
- [100 Days of Kubernetes](#)

## Experiment with side projects

- Trying examples is a great way of learning
- Just play around and experiment with workloads
- Celebrate small successes

## Join Kubernetes communities

- [Kubernetes Slack](#)
- [Cloud Native Linz](#)
- [Cloud Native Austria](#) (Graz, Innsbruck, Salzburg & Vienna)



Thank you! :)

Happy to take any  
questions.