

Workshop “Kubernetes”




AT Computing
(Vijfhart Group)



Introduction containers (1)

Conventional production environment

- all applications use same 'ecosystem'
 - filesystem
 - additional storage
 - network
 - process numbers (PIDs)
 - user identities
 - hardware resources
- disadvantages
 - laborious (de)installation of application
 - compromised application might
 - access all data
 - manipulate entire network stack
 - compromise other applications
 - application might overload resources



Introduction containers (2)

Container: isolated ecosystem for application

- image with own mini-filesystem containing
 - programs
 - configuration files
 - data files
 - ...


- own additional storage

- own network

- own PID numbers

- own user identities

- cpu/memory/disk limitations



Introduction containers (3)

Image

- needed to start container on destination host
- contains
 - mini-filesystem
 - metadata
- *developer* ('dev')
 - builds image for application
 - ships image to registry
- *operations* ('ops')
 - pulls image from registry
 - uses image to activate container
 - in any environment: test, acceptance, production
 - on any operating system: Linux, Windows, macOS
 - on any platform: physical host, virtual machine, cloud

Workshop “Kubernetes”




Container summary using Docker



Container summary

Docker container




- run application in isolated environment
 - own mini- filesystem
 - own network
 - own private PID numbers
 - own mounted filesystems
 - own users
 - separate root privileges
 - cpu/memory limitation
- lightweight
- simple
- large community (lots of images)

Container summary – virtual machines vs. containers


Virtual machine

- application
- libraries
- full command set
- full operating system (kernel)




Virtual machines

Docker container



- application
- libraries
- limited command set
- no private kernel
- *container process is native process for host!*



Docker containers

Container summary – run from base image

Start container from base image

- example: run command in base image **ubuntu** from Docker registry

```
$ docker run image ubuntu overruling command ps -f
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	08:59	?	00:00:00	ps -f

```
$
```

- container terminates when process in container terminates
- example: run command in base image **ubuntu** interactively

```
$ docker run interactive tty -it ubuntu bash
```

```
root@f8d91cbaca03:/# ps -f
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	1	09:12	?	00:00:00	bash
root	11	1	0	09:12	?	00:00:00	ps -f

```
root@f8d91cbaca03:/# exit
```

```
$
```


Container summary – build custom image

Build custom image

- specify own modifications in file **Dockerfile**

```
$ cat Dockerfile
FROM ubuntu:18.04
RUN apt-get update && apt-get install -y apache2
COPY index.html /var/www/html/index.html
CMD ["/usr/sbin/apache2ctl", "-D", "FOREGROUND"]
```

```
$ cat index.html
<h1> Message from container! </h1>
```

- build custom image

```
$ docker build -t atcomp/apachetest .
Successfully built 5d3b567581df
```

new image

directory containing **Dockerfile**
and other files needed in image

- list images

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
atcomp/apachetest	latest	5d3b567581df	About an hour ago	268.1 MB
docker.io/ubuntu	18.04	104bec311bcd	3 months ago	128.9 MB

Container summary – run from custom image

Start container from custom image

- run custom container

```
$ docker run -p 8080:80 -d atcomp/apachetest  
c8eda1f3a6734304195ab3e47280ee77c719fa5c365ba428bc2037e250754c2e
```

SHA256 (often abbreviated with first 48 bits)

- contact webserver via URL `http://localhost:8080` (web browser/curl)

```
$ curl http://localhost:8080  
<h1> Message from container! </h1>
```

- show running containers

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	PORTS
c8eda1f3a673	atcomp/apachetest	"/usr/sbin/apache2ctl"	0.0.0.0:8080->80/tcp

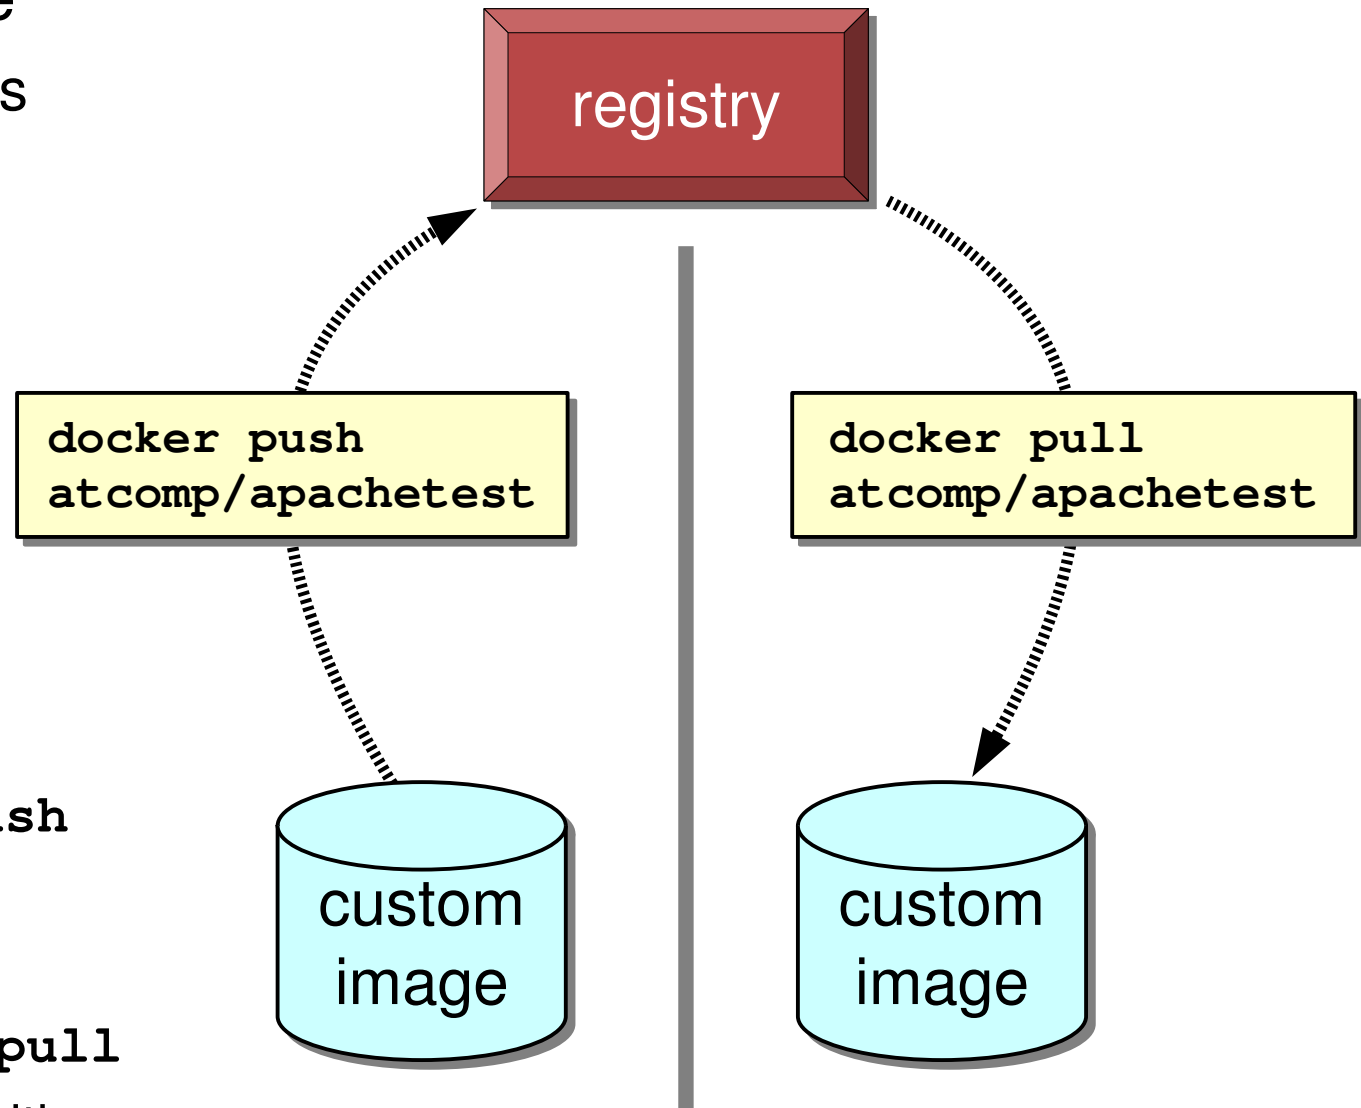
- terminate running container

```
$ docker stop c8eda1f3a673  
c8eda1f3a673
```

Container summary – registries

Docker registry: image store

- server(s) containing images
 - multiple versions
- possibilities
 - Docker Hub (100,000+)
 - in-company registry
- images can be
 - *pushed*
 - stored with `docker push`
 - *pulled*
 - explicitly with `docker pull`
 - implicitly on initial use with `docker run`
 - implicitly on initial use with **FROM** in **Dockerfile**



Workshop “Kubernetes”



Introduction



What is Kubernetes?

Kubernetes ('helmsman') a.k.a. K8s

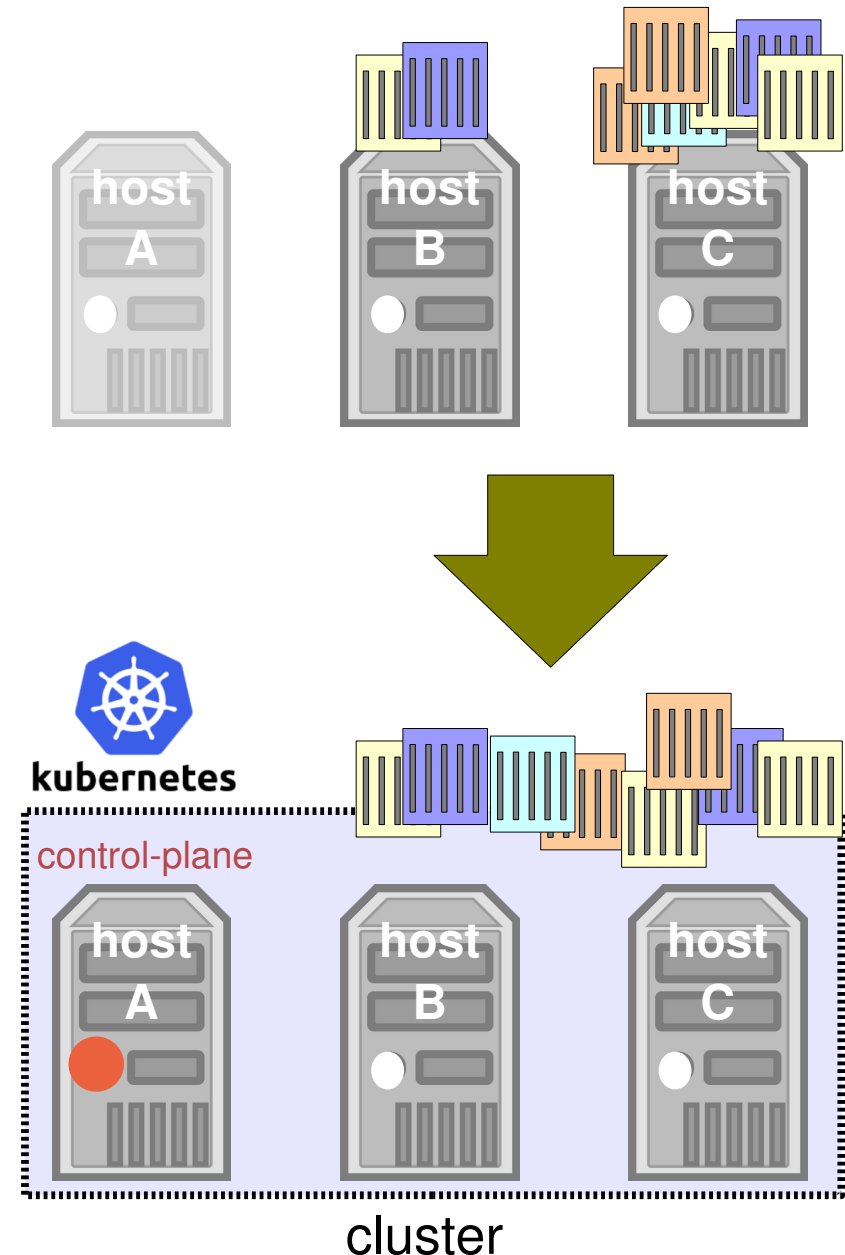
- combines various hosts into cluster
 - scalability
 - reliability (failover)
- orchestrates containers
 - activate
 - monitor
 - terminate
- can manage various container implementations, like *containerd*, *CRI-O*, *Docker*,



Kubernetes orchestration

Kubernetes orchestration

- introduced in 2014, inspired by Google's Borg
- maintained by *Cloud Native Computing Foundation (CNCF)*
- open source
- concept
 - cluster needs at least one *control-plane node*, formerly called *master node*
 - other hosts in cluster are *worker nodes*
 - only run container instances

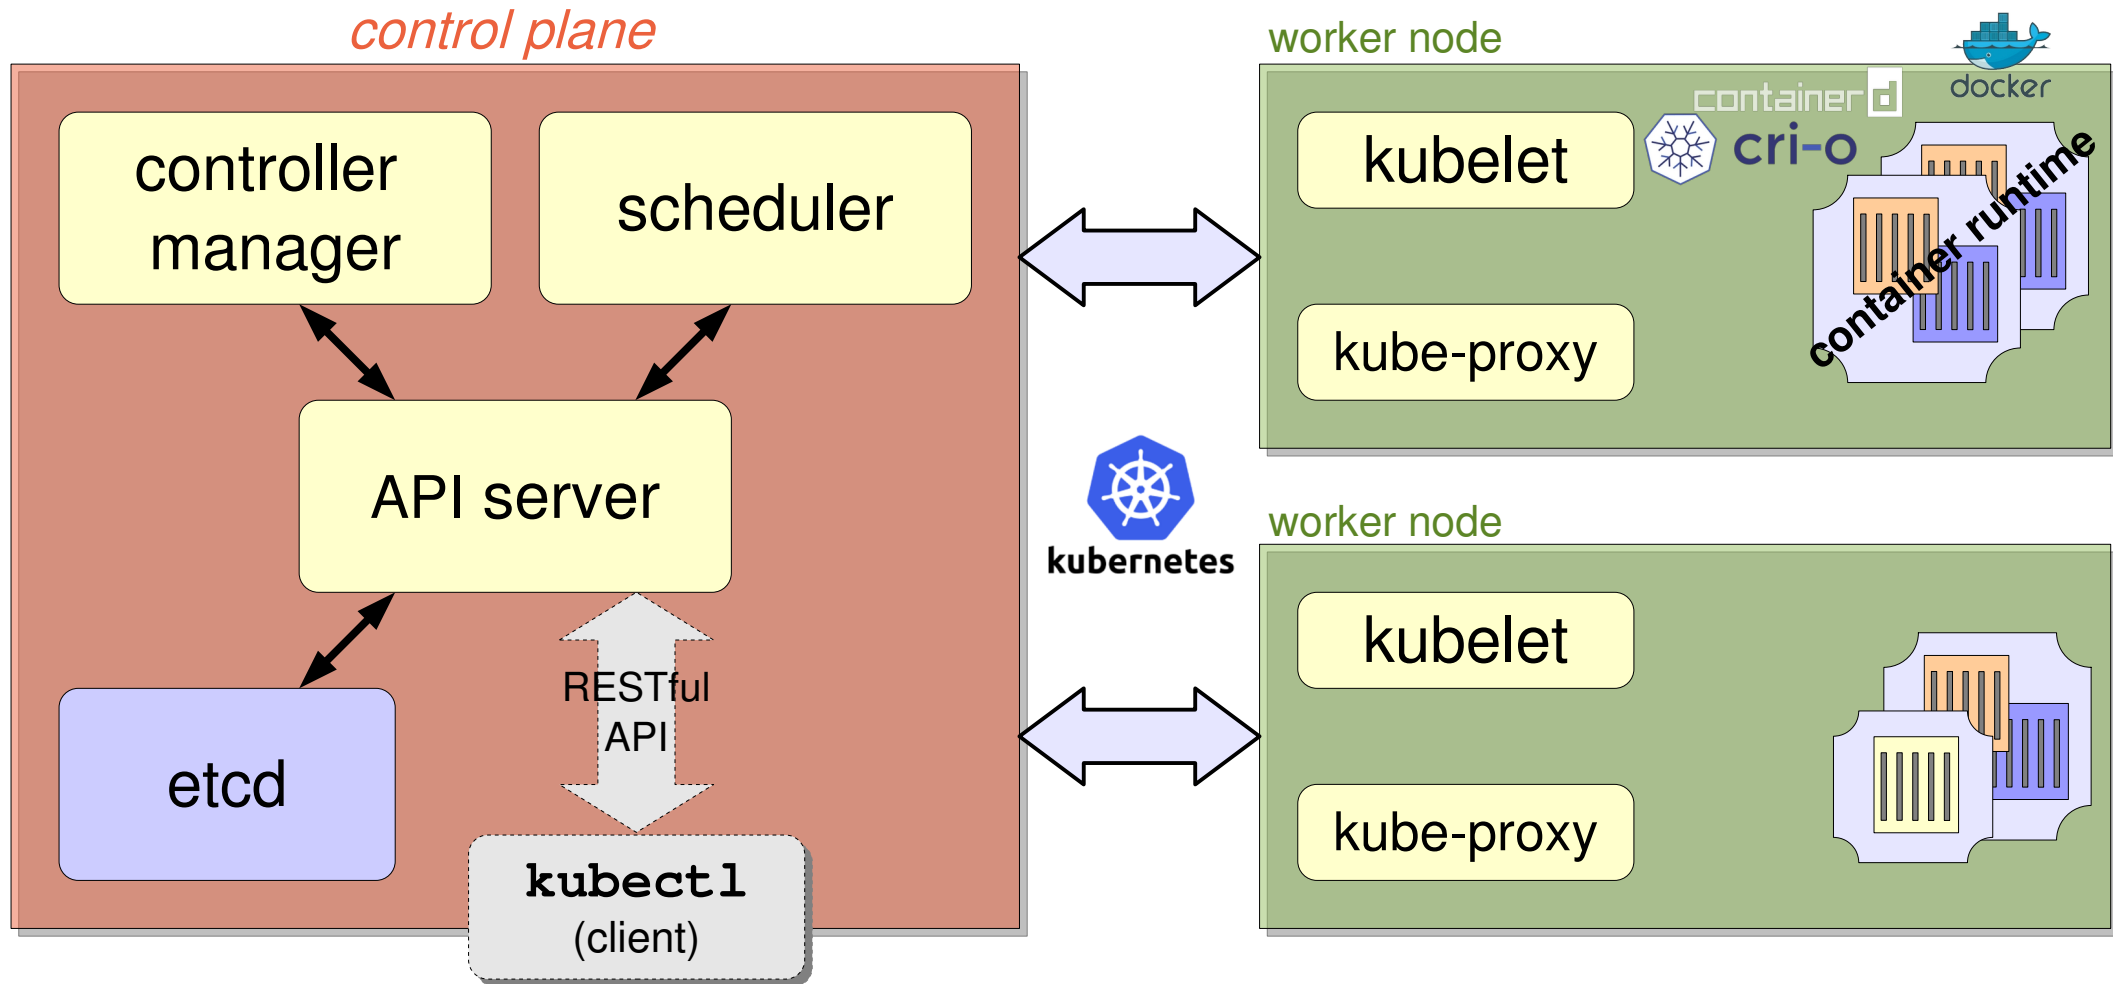


User interface

User interfaces

- command line interface: `kubectl subcommand object [options]`
 - subcommands
 - create, modify and delete objects: `create, apply, delete`
 - query current state of objects: `get, describe`
 - other actions, like: `logs, exec,`
 - overview of subcommands: `kubectl`
 - more info about subcommand: `kubectl subcommand --help`
- graphical user interface
 - also valid for cloud implementations, like
 - Google Kubernetes Engine (GKE)
 - Amazon Elastic Kubernetes Service (EKS)
 - Azure Kubernetes Service (AKS)

Kubernetes architecture





- *API server*: controls entire cluster
- *scheduler*: schedules containers (via pods) on nodes
- *controller manager*: controls required number of replicas
- *etcd*: distributed key-value store to maintain current cluster state

- *kubelet*:
 - pod startup & monitoring
 - resource management
- *kube-proxy*: container exposure to network & load balancing

Pods

Pod

- conceptually smallest unit, running one application
- consists of one or more closely related containers
 - typically one container running primary application
 - additional containers supporting primary application (if needed)
- containers in same pod
 - run on same worker node
 - share *storage volumes*
 - allows persistent data
 - allows containers to access same storage
 - share *network*
 - same IP address on pod network, same ports
 - allows containers to communicate via **localhost**



Pods – startup, status and termination

Pod lifetime

- activate pod specifying command line parameters

```
$ kubectl run apatest --restart=Never --image=atcomp/apachetest --port=80  
pod/apatest created
```

or

- activate pod via manifest file

```
$ kubectl apply -f apa.yml  
pod/apatest created
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: apatest  
spec:  
  containers:  
  - name: apacont  
    image: atcomp/apachetest  
    ports:  
    - containerPort: 80  
  restartPolicy: Never
```

- verify pod status and delete pod

```
$ kubectl get pods  
NAME      READY   STATUS    RESTARTS   AGE  
apatest  1/1     Running   0           82s  
  
$ kubectl get pods -o wide  
NAME      READY   STATUS    RESTARTS   AGE   IP            NODE  
apatest  1/1     Running   0           83s   10.244.1.254  hostb  
  
$ kubectl delete pod/apatest  
pod "apatest" deleted
```

Pod completion

After container termination

- pod is 'completed' but not removed

```
$ kubectl run testpod --restart=Never --image=ubuntu -- sleep 10
pod/testpod created
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
testpod	1/1	Running	0	6s

```
$ kubectl get pods
```

after a while...

NAME	READY	STATUS	RESTARTS	AGE
testpod	0/1	Completed	0	17s

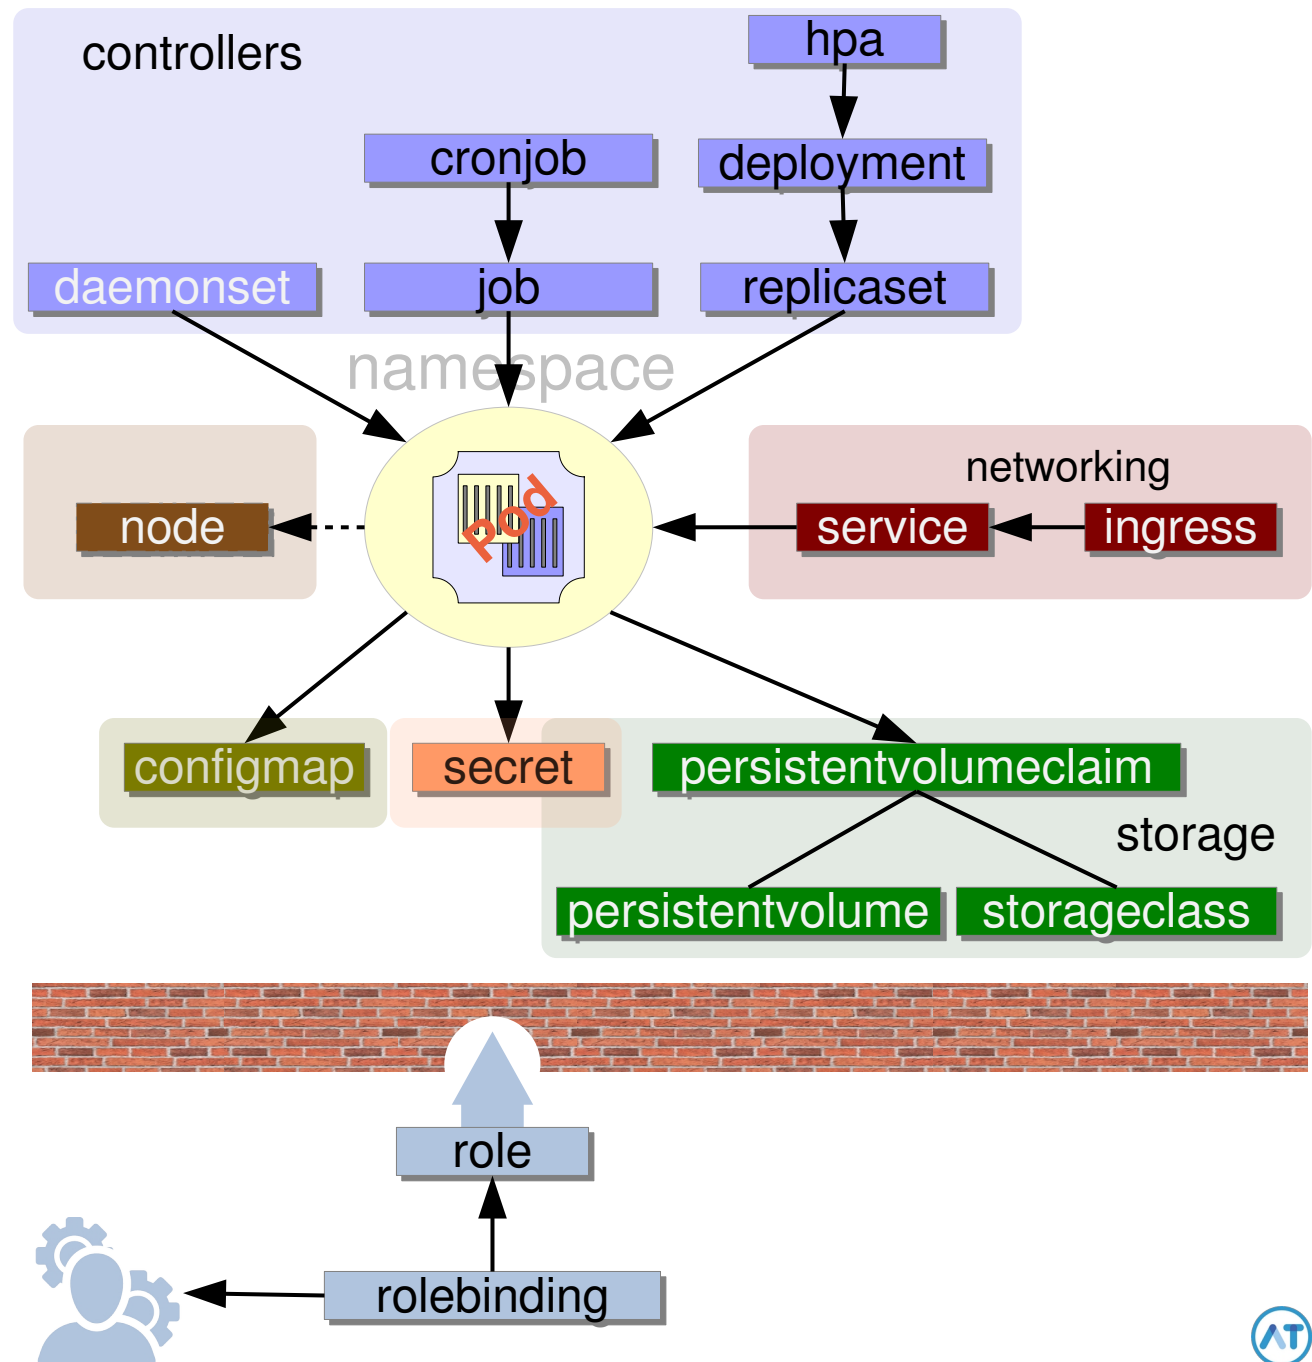
- logs can still be watched
- events can still be verified
- pods are not 'self-healing'
 - *controller* needed to restart completed pod (covered later)
- interactive pod (`-it`)

```
$ kubectl run testpod -it --restart=Never --image=ubuntu -- bash
```

Kubernetes objects (1)

Kubernetes concept

- numerous object types
- every object has
 - type ('kind')
 - unique name
- object refers to other objects by using
 - labels
(most references *to* pods)
 - or
 - object names
(most references *from* pods)



Kubernetes objects (2)

Kubernetes object

- has *desired state* and *actual state* (current state)
 - control plane continuously tries to match *actual state* to *desired state*
 - desired state defined by **spec** in manifest file

- manifest consists of

apiVersion:	version of API
kind:	type of object
metadata:	object name (must be unique) optional labels for selection
spec:	specification of desired state, depending on type

```
apiVersion: v1
kind: Pod
metadata:
  name: apatest
  labels:
    app: webserver
spec:
  containers:
  - name: apacont
    image: atcomp/apachetest
    ports:
    - containerPort: 80
  restartPolicy: Never
```

apa.yml

Labels

Object labels

- every object has unique name
- additionally, *labels* can be assigned to be used for
 - selection on command line with `-l` flag

```
$ kubectl get pods -l app=webserver
$ kubectl delete all -l app=webserver
```

- refer from one object to another object
e.g. to assign **Service** object to **Pod** object

apa.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: apatest
  labels:
    app: webserver
spec:
  ....
```

apa-svc.yml

```
apiVersion: v1
kind: Service
metadata:
  name: webservice
spec:
  ports:
    - port: 80
    selector:
      app: webserver
```

Namespaces

Namespaces: subdivide cluster into various virtual clusters

- per project, per application, per developer team, per department, per

```
$ kubectl create ns devel
namespace/devel created
```

```
$ kubectl get ns
NAME          STATUS    AGE
default       Active    278d
devel         Active    14s
....
```

```
$ kubectl apply -f apa.yaml -n devel
```

in namespace *devel*

```
$ kubectl get pods -n devel
```

in namespace *devel*

NAME	READY	STATUS	RESTARTS	AGE
apatest	1/1	Running	0	48s

- allow
 - separate scope for object names
 - limitation on resource utilization (cpu, memory, storage, number of pods, ...)

Workshop “Kubernetes”




Controllers



Controllers – an introduction

Controllers

- pod ('naked pod', 'bare pod') is not self-healing
- pods usually started under supervision of *controller*
 - create, manage, monitor and restart pods
 - provide rolling updates
- part of *controller manager* in control plane
- manifest file of controller requires
 - controller-specific definitions
 - template of **pod** (to be controlled)



```
apiVersion: apps/v1
kind: somecontroller
metadata:
  name: sleeper
  ....
spec:
  controller-specific stuff...
  template:
    metadata:
      labels:
        app: webserver
    spec:
      containers:
        - name: snorecont
          image: ubuntu
          command: ["sleep", "60"]
```

pod template

Controller types

Controller types

- *ReplicaSet* (rs) – preferably combined with Deployment
 - pod replicas and restart of failing pods
- *Deployment* – recommended
 - pod replicas by using *ReplicaSet*
 - rolling updates and rollbacks

Deployment – create

Deployment object

- definition implies **ReplicaSet** and **Pod**
- create deployment

```
$ kubectl apply -f apa-deploy.yml
deployment.apps/apadep created
```

```
$ kubectl get deployment
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	...
apadep	3	3	3	3	

```
$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
apadep-7b6fc56c77	3	3	3	18s

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	...
apadep-7b6fc56c77-6ldcv	1/1	Running	0	
apadep-7b6fc56c77-bqxhr	1/1	Running	0	
apadep-7b6fc56c77-hk7qp	1/1	Running	0	


```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apadep
  labels:
    app: webserver
spec:
  replicas: 3
  selector:
    matchLabels:
      app: apapod
  template:
    metadata:
      labels:
        app: apapod
    spec:
      containers:
        - name: apacont
          image: atcomp/apachetest:1.14
          ports:
            - containerPort: 80
```

apa-deploy.yml

pod template

Deployment
image: apachetest:1.14

ReplicaSet
replicas: 3
image: apachetest:1.14



Deployment – rolling update

Rolling update

- executed when pod specification changes, like new image

```
$ kubectl set image deployment/apadep apacont=atcomp/apachetest:1.15
deployment.apps/apadep image updated
```

```
$ kubectl get all -o wide
```

NAME	READY	UP-TO-DATE	AVAILABLE	CONTAINERS	IMAGES	...
deployment.apps/apadep	3/3	3	3	apacont	atcomp/apachetest:1.15	


NAME	DESIRED	CURRENT	READY	CONTAINERS	IMAGES	...
replicaset/apadep-5b4f756b5c	3	3	3	apacont	atcomp/apachetest:1.15	
replicaset/apadep-7b6fc56c77	0	0	0	apacont	atcomp/apachetest:1.14	

NAME	READY	STATUS	IP	...
pod/apadep-5b4f756b5c-5kvrnx	1/1	Running	10.244.2.213	
pod/apadep-5b4f756b5c-rflpn	1/1	Running	10.244.2.212	
pod/apadep-5b4f756b5c-z5zbj	1/1	Running	10.244.1.89	



- creates new replicaset within deployment
 - pod replicas replaced one-by-one
 - original replicaset preserved for rollback

```
$ kubectl rollout undo \
    deployment/apadep
```



Deployment – scaling

Deployment – scaling

- modify number of replicas

```
# kubectl scale --replicas=2 deployment/apadep  
deployment.extensions/apadep scaled
```

```
$ kubectl get all -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
pod/apadep-5b4f756b5c-5kvrnx	1/1	Running	3	4m20s	10.244.2.213	hostc
pod/apadep-5b4f756b5c-rflpn	1/1	Running	3	4m19s	10.244.1.212	hostb

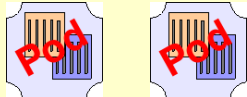
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	CONTAINERS	IMAGES
deployment.apps/apadep	2	2	2	2	apacont	atcomp/apachetest:1.15

NAME	DESIRED	CURRENT	READY	CONTAINERS	IMAGES
replicaset.apps/apadep-5b4f756b5c	2	2	2	apacont	atcomp/apachetest:1.15
replicaset.apps/apadep-7b6fc56c77	0	0	0	apacont	atcomp/apachetest:1.14

- replicas can even be scaled to 0 (temporarily no pods)
- autoscaling possible with HPA controller
 - minimum and maximum number of replicas
 - target CPU time utilization per replica

Deployment
image: apachetest:1.15

ReplicaSet
replicas: 2
image: apachetest:1.15



Workshop “Kubernetes”




Networking



Kubernetes networking

Communication possibilities

- container-to-container in same pod
 - via `localhost` (loopback interface)
 - port numbers may conflict between applications
- pod-to-external
 - via Network Address Translation (NAT)
- pod-to-pod
 - what is IP address of destination pod?
- external-to-pod
 - what is IP address of destination pod?



Services – IP address of pod

Example: access via dynamic IP address

- create deployment with 2 pod replicas

```
$ kubectl apply -f apa-deploy.yml
deployment.apps/apadep created

$ kubectl get pods -o wide
NAME                                READY ... IP                NODE
apadep-8654d77c94-928tw            1/1      10.244.1.130             hostb
apadep-8654d77c94-ggv8c            1/1      10.244.2.64              hostc

$ curl 10.244.1.130
<h1> Message from container! </h1>
```

- disadvantages
 - no access from outside cluster
 - no load balancing
 - when pod terminates, it probably gets another IP address after restart


```
$ kubectl delete pod/apadep-....-928tw
$ kubectl get pods -o wide
NAME                                READY IP                NODE
apadep-8654d77c94-7br4b            1/1   10.244.1.131         hostb
apadep-8654d77c94-ggv8c            1/1   10.244.2.64          hostc
```

```
apa-deploy.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apadep
  labels:
    app: webserver
spec:
  replicas: 2
  selector:
    matchLabels:
      app: apapod
  template:
    metadata:
      labels:
        app: apapod
    spec:
      containers:
        - name: apacont
          image: atcomp/apachetest:1.14
          ports:
            - containerPort: 80
```


Services – internal access

Service object

- gets static virtual IP (VIP) address, though dynamically assigned
 - stable IP address to reach mortal pods
 - load balancing
- contains selector
 - refers to label of pod to attach to
 - endpoint object created per target port
- accessibility determined by *type*
 - ClusterIP**: routable within cluster (default) for internal access
 - NodePort**: static port on every node in cluster for external access



Services – internal: setup ClusterIP

Example: add *internal* service for webserver

- create service referring to label **app=apapod**
- type **ClusterIP** to provide internal access

```
$ kubectl apply -f apa-svc.yml
```

service/webserv created

```
$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	PORT(S)
webserv	ClusterIP	10.97.15.1		80/TCP	

```
$ kubectl get ep
```

NAME	ENDPOINTS	AGE
webserv	10.244.1.131:80, 10.244.2.64:80	67m

```
$ kubectl get all
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/apadep	2	2	2	2	10m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/apadep-8654d77c94	2	2	2	10m

NAME	READY	STATUS	RESTARTS	AGE
pod/apadep-8654d77c94-7br4b	1/1	Running	0	10m
pod/apadep-8654d77c94-ggv8c	1/1	Running	0	10m

apa-deploy.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apadep
spec:
  replicas: 2
  selector:
    ....
  template:
    metadata:
      labels:
        app: apapod
    spec:
      ....
```


apa-svc.yml

```
apiVersion: v1
kind: Service
metadata:
  name: webserv
  labels:
    app: apa
spec:
  type: ClusterIP
  ports:
    - port: 80
      targetPort: 80
      protocol: TCP
  selector:
    app: apapod
```

Services – internal: discovery

Service discovery by other pods

- via internal DNS
- maintains record for every service:
`service[.ns.svc.cluster.local]`
- preferred (always available)!



```
$ kubectl run mypod -it --restart=Never --image=atcomp/nwubuntu
root@mypod:/# cat /etc/resolv.conf
nameserver 10.96.0.10
search default.svc.cluster.local svc.cluster.local cluster.local


root@mypod:/# host webserv
webserv.default.svc.cluster.local has address 10.97.15.1

root@mypod:/# curl webserv
<h1> Message from container! </h1>
```

Services – external access

External access to pod

- requires service type **Nodeport**
(implies service type **ClusterIP**)
- port range 30000-32767
 - specify with keyword **nodePort**
 - dynamically assigned when keyword **nodePort** omitted



Services – external: setup NodePort

Example: add *external* service for webserver

- create service

```
$ kubectl apply -f apa-svcn.yml
service/webserv created
```

```
$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	PORT(S)
webserv	NodePort	10.97.15.1		80:32123/TCP	

```
$ ss -tl
```

State	Recv-Q	Send-Q	Local Address	Foreign Address
LISTEN	0	0	:::32123	:::*

- access from any host outside cluster

```
anyhost$ curl hosta:32123           or hostb or hostc
<h1> Message from container! </h1>
```

- access from inside cluster (similar to ClusterIP)

```
$ kubectl run mypod -it --restart=Never
                        --image=atcomp/nwubuntu
root@mypod:/# curl webserv
<h1> Message from container! </h1>
```

apa-deploy.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apadep
  labels:
    app: apa
spec:
  replicas: 2
  selector:
    ....
  template:
    metadata:
      labels:
        app: apa
    spec:
      containers:
        ....
```

apa-svcn.yml

```
apiVersion: v1
kind: Service
metadata:
  name: webserv
  labels:
    app: apa
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 32123
      protocol: TCP
  selector:
    app: apa
```

Workshop “Kubernetes”




Storage



Kubernetes Volumes

Storage

- *stateless pod* – preferred!
 - no need to preserve data
 - container storage (filesystem) on local disk
- *stateful pod*
 - requires persistent volume
 - volume can be mapped on
 - local disk
 - restricted use
 - network storage
 - pod terminated on **hostB** might be restarted on **hostC** and/or
 - pods running on different hosts might share same storage



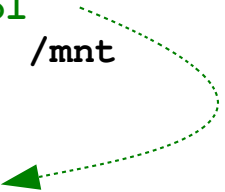
Volumes

Volumes

- have specific type
- have explicit lifetime
(life span of pod, permanent, ...)
- in pod manifest
 - spec.volumes:**
provided volumes of certain type
 - spec.containers.volumeMounts:**
mount point in container

```
apiVersion: v1
kind: Pod
metadata:
  name: ....
spec:
  containers:
    - name: ....
      image: ....
      volumeMounts:
        - name: myvol
          mountPath: /mnt

  volumes:
    - name: myvol
      volumeType
```



Volumes – shared between pods on same node

Persistent shared volume on node

- persistent storage in filesystem of host
 - notice: only host on which pod is created!
 - read/write access

solution: **hostPath** volume

- example: run pod on specific host

```
$ kubectl get nodes --show-labels
NAME      STATUS   ROLES    ... LABELS
hosta     Ready    master   kubernetes.io/hostname=hosta
hostb     Ready    <none>    kubernetes.io/hostname=hostb
hostc     Ready    <none>    kubernetes.io/hostname=hostc


$ kubectl apply -f hostpath-pod.yml

$ kubectl get pod/apahost -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE
apahost   1/1     Running   0           19m   10.244.2.240 hostc

$ curl 10.244.2.240
<H1> Welcome to hostc! </H1>
```

hostpath-pod.yml

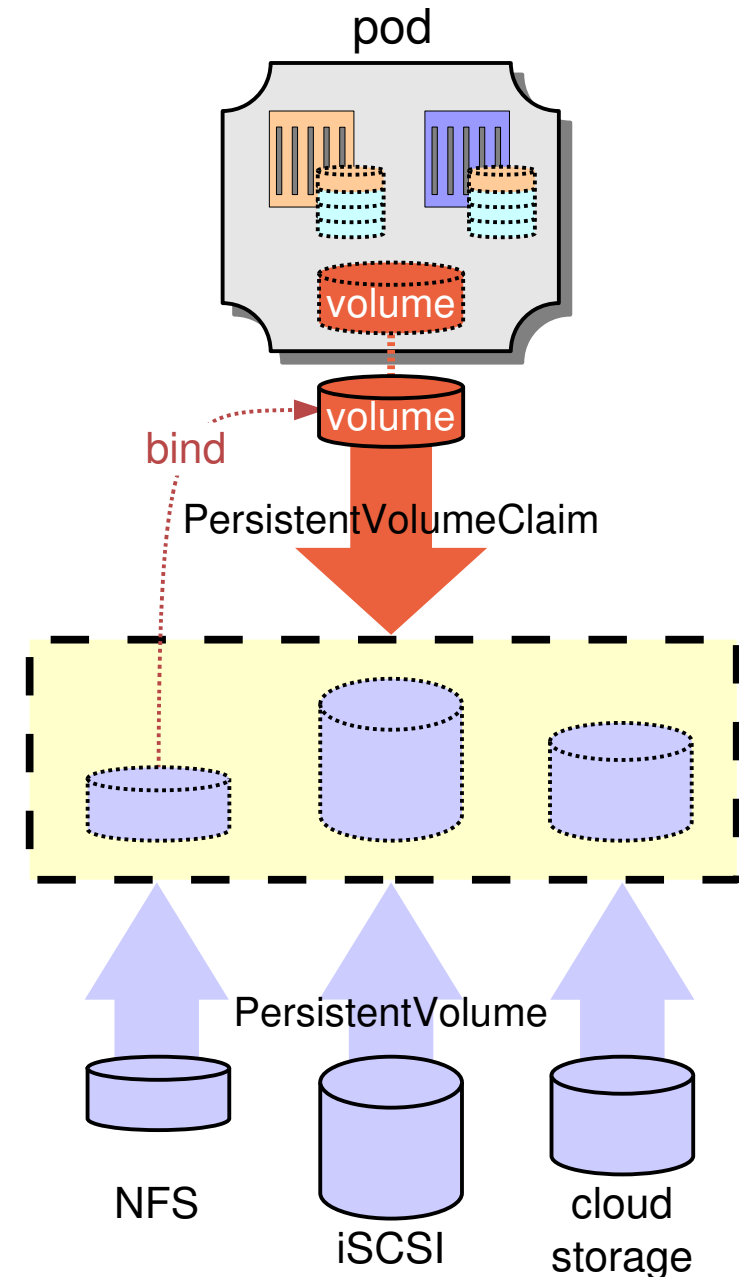
```
apiVersion: v1
kind: Pod
metadata:
  name: apahost
  labels:
    app: webserver
spec:
  nodeSelector:
    kubernetes.io/hostname: hostc
  containers:
  - name: apatest
    image: atcomp/apachetest
    ports:
    - containerPort: 80
  volumeMounts:
  - name: hosthtml
    mountPath: /var/www/html
  volumes:
  - name: hosthtml
    hostPath:
      path: /var/www/html
      type: Directory
```



Persistent Volumes – static provisioning (1)

Persistent volumes – *static allocation*

- managed by
 - **PersistentVolume** (PV)
 - *piece of storage* provisioned by administrator
 - example types: NFS, iSCSI
 - **PersistentVolumeClaim** (PVC)
 - *request for storage* by user to be mounted in pod
 - specific properties can be defined, like size, access mode, performance,
- binding of PVC to PV
 - 1-to-1 mapping
 - PVC state '**Pending**' if no suitable PV available



Persistent Volumes – static provisioning (2)

Example persistent volume

- create PV of 1GiB based on NFS
- request PVC of 500MiB

```
$ kubectl apply -f pub-pv.yml
persistentvolume/pub-pv created
```

```
$ kubectl get pv
```

NAME	CAPACITY	ACCESS ...	STATUS	CLAIM
pub-pv	1Gi	RWX	Available	

```
$ kubectl apply -f pub-pvc.yml
persistentvolumeclaim/pub-pvc created
```

```
$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS...
pub-pvc	Bound	pub-pv	1Gi	RWX

```
$ kubectl get pv
```

NAME	CAPACITY	ACCESS...	STATUS	CLAIM
pub-pv	1Gi	RWX	Bound	default/pub-pvc

pub-pv.yml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pub-pv
spec:
  storageClassName: nfspool
  accessModes:
    - ReadWriteMany
  capacity:
    storage: 1Gi
  nfs:
    server: nasi
    path: /nfs/Public
    readOnly: false
```

pub-pvc.yml

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pub-pvc
spec:
  storageClassName: nfspool
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 500Mi
```

Persistent Volumes – static provisioning (3)

Example persistent volume – cont'd

- create deployment with pod using PVC as volume
- pod pending as long as PVC pending

```
$ kubectl apply -f pub-deploy.yml
deployment/pubdep created

$ kubectl get pod
NAME                                READY   STATUS ...
pubdep-5874f6fbd6-qvndg            1/1     Running

$ kubectl exec -it pubdep-...-qvndg bash
root@pubdep-5874f6fbd6-qvndg:/# ls -l /public
....
drwxrwxrwx+ ... 4096 Jun 22 2018 Documents
drwxrwxrwx+ ... 4096 Feb 11 2016 Music
drwxrwxrwx+ ... 4096 Dec 17 2016 Photos
```

```
pub-deploy.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: pubdep
  labels:
    app: pubdep
spec:
  replicas: 1
  selector:
    matchLabels:
      app: pubpod
  template:
    metadata:
      labels:
        app: pubpod
    spec:
      containers:
        - name: sleepcont
          image: ubuntu
          command: ["sleep", "3600"]
          volumeMounts:
            - name: pubstore
              mountPath: /public
      volumes:
        - name: pubstore
          persistentVolumeClaim:
            claimName: pub-pvc
```

Persistent Volumes – dynamic provisioning

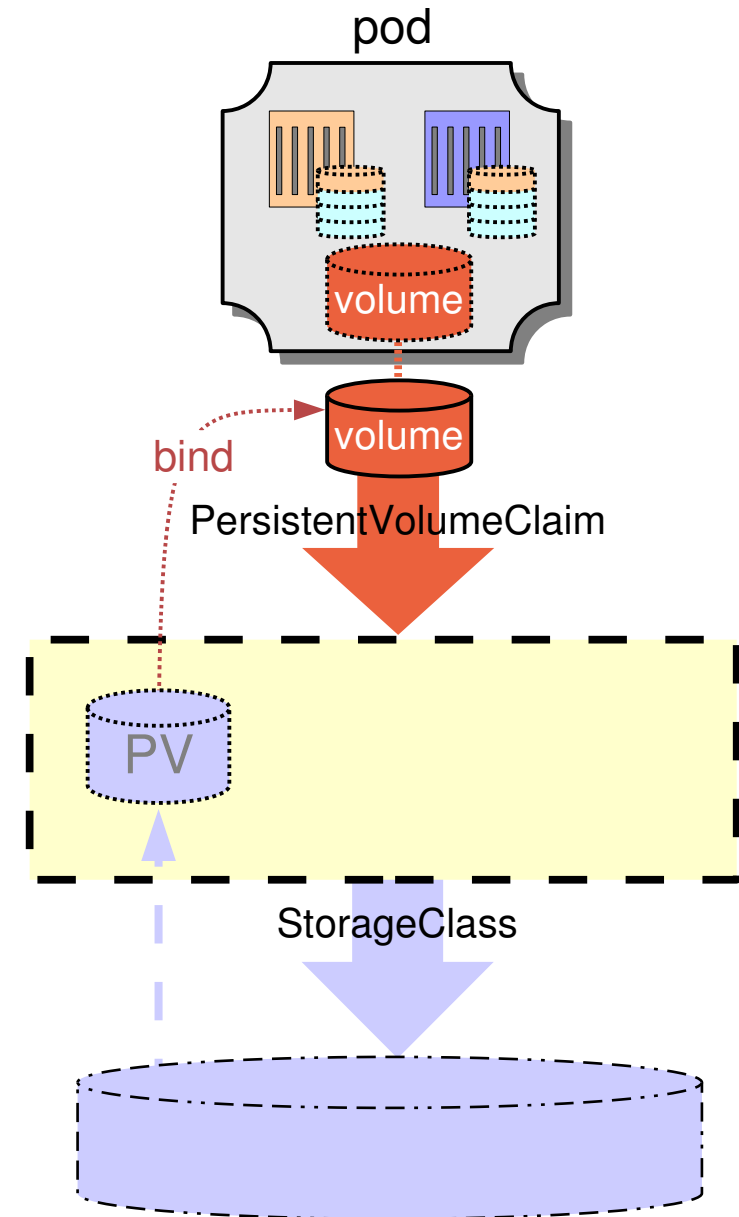
Persistent volumes – *dynamic provisioning*

- managed by **StorageClass** (SC)
 - dynamically allocates storage when claim (PVC) issued
 - uses *provisioner* (volume plugin) for allocation

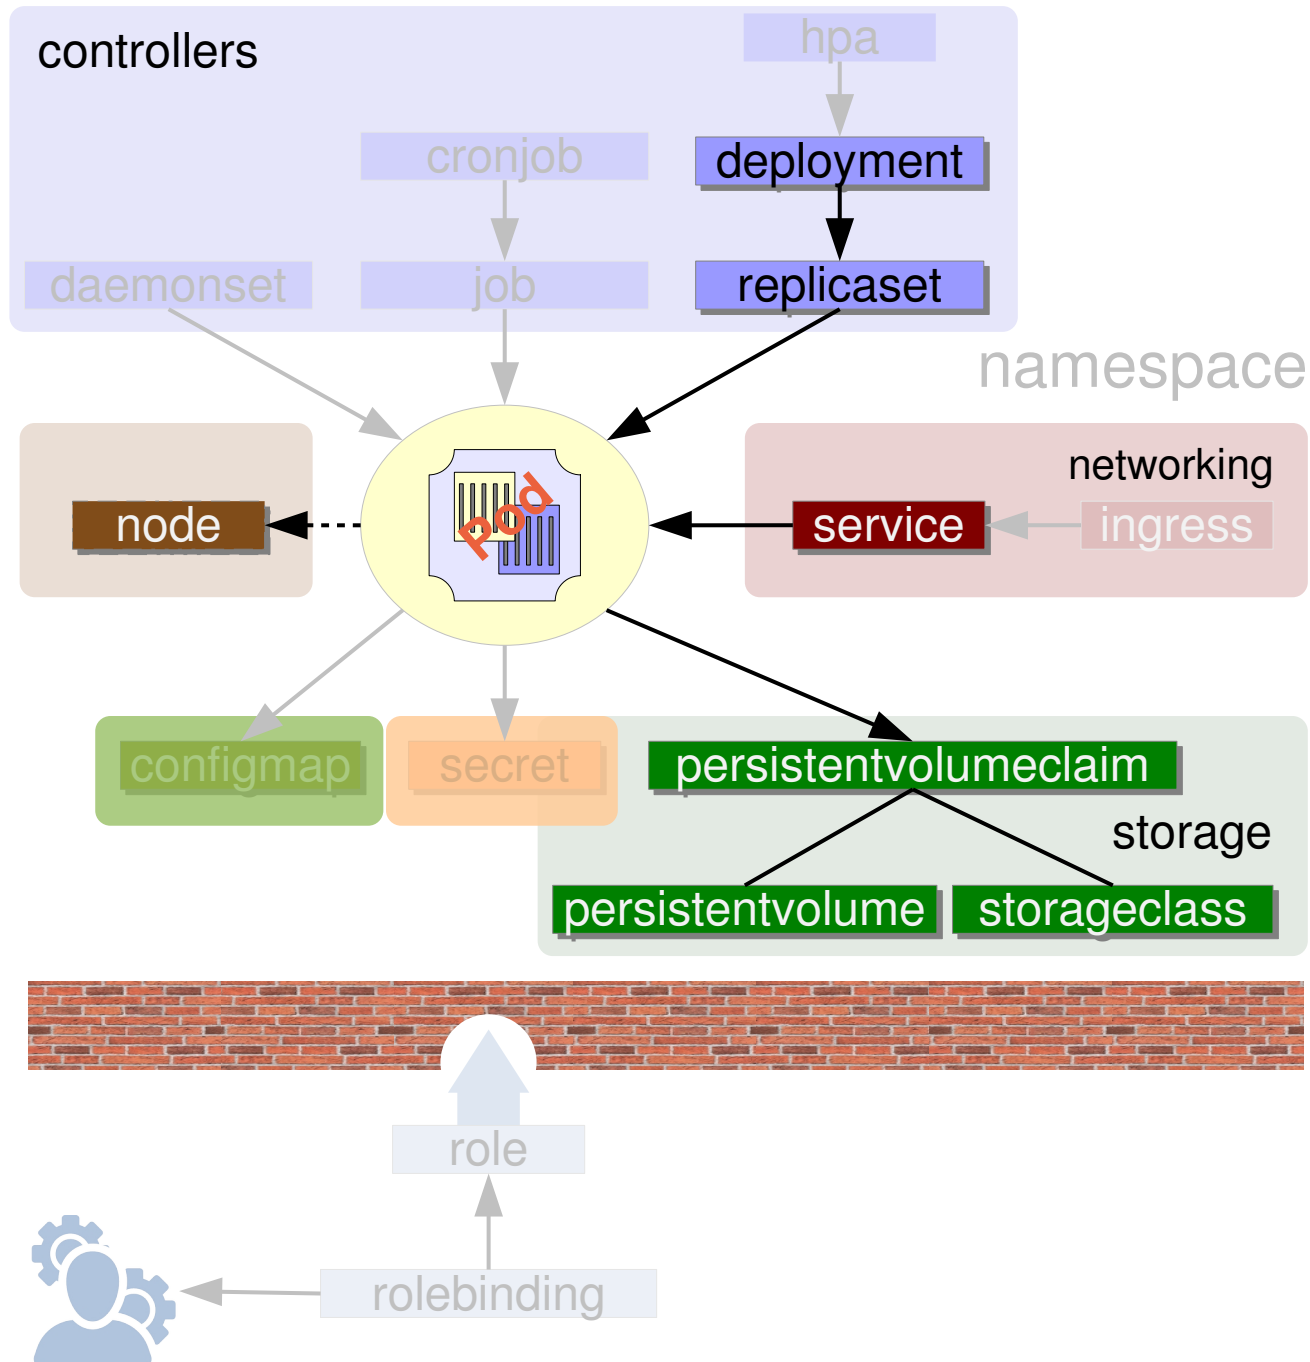
builtin provisioners (**kubernetes.io**):

gce-pd	(GCEPersistentDisk)
aws-ebs	(AWSElasticBlockStore)
azure-disk	(AzureDisk)
azure-file	(AzureFile)
glusterfs	(Glusterfs)
and many others....	

- dynamically creates PV object
- volume initially empty



Objects in this workshop



Workshop “Kubernetes”



Workshop is extraction from the course
“Kubernetes Fundamentals”
(three days)

Questions?